

Database Systems

Chapter 2: Data Models

- Data Models defined.
- Look @ some Data Models
- Look @ Relational Data Model
- Look @ SQL in MySQL

What's a Data Model

- Structure of the Data
 - Tree, Graph
 - Semi-Structured (XML)
 - Probabilistic Graphs
 - Relational
- Operations on the Data
- Constraints on the Data

What's do we get from Data Model

- Efficiency!
 - Provide us an abstraction for data.
 - Operations within abstractions optimized.
- Mathematical formalisms for proofs.

Tree Model File System

- Structure of the Data

- Tree

- Operations on the Data

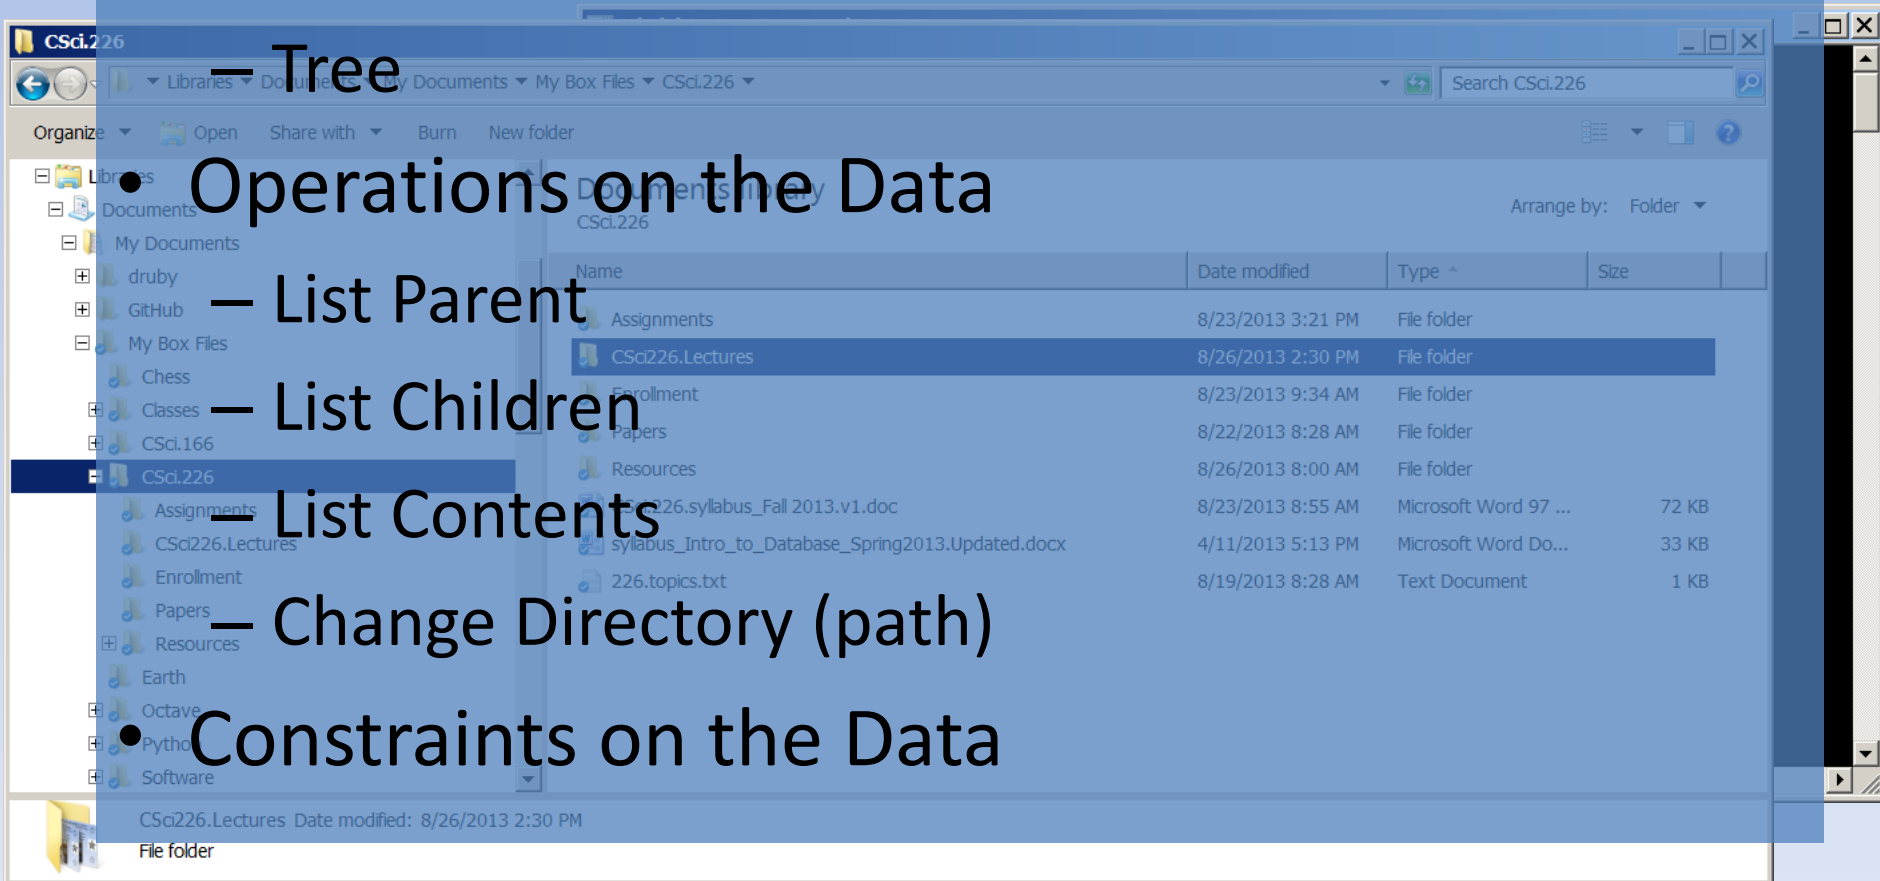
- List Parent

- List Children

- List Contents

- Change Directory (path)

- Constraints on the Data



Semi-Structured Data (XML)

<Bookstore>

<Book ISBN="ISBN-0-13-713526-2" Price="85" Edition="3rd">

<Title>A First Course in Database Systems </Title>

<Authors>

<Auth> <FName>Jeffrey</FName> <LName>Ullman</LName> </Auth>

<Auth><FName>Jennifer</FName><LName>Widom</LName></Auth>

</Authors>

</Book>

<Book ISBN="ISBN-0-13-815504-6" Price="100">

<Remark> Buy this book bundled with "A First Course" - a great deal!</Remark>

<Title>Database Systems: The Complete Book</Title>

<Authors>

<Auth><FName>Hector</First_Name><LName>Garcia-Molina</LName></Auth>

<Auth><FName>Jeffrey</FName><LName>Ullman</LName></Auth>

<Auth><FName>Jennifer</FName><LName>Widom</LName></Auth>

</Authors>

</Book>

</Bookstore>

XML

- Semantic Tags:
 - `<Bookstore> elements </Bookstore>`
- Attributes
 - `<Book ISBN="ISBN-0-13-713526-2" Price="85" Edition="3rd">
</Book>`
- Nested Elements:

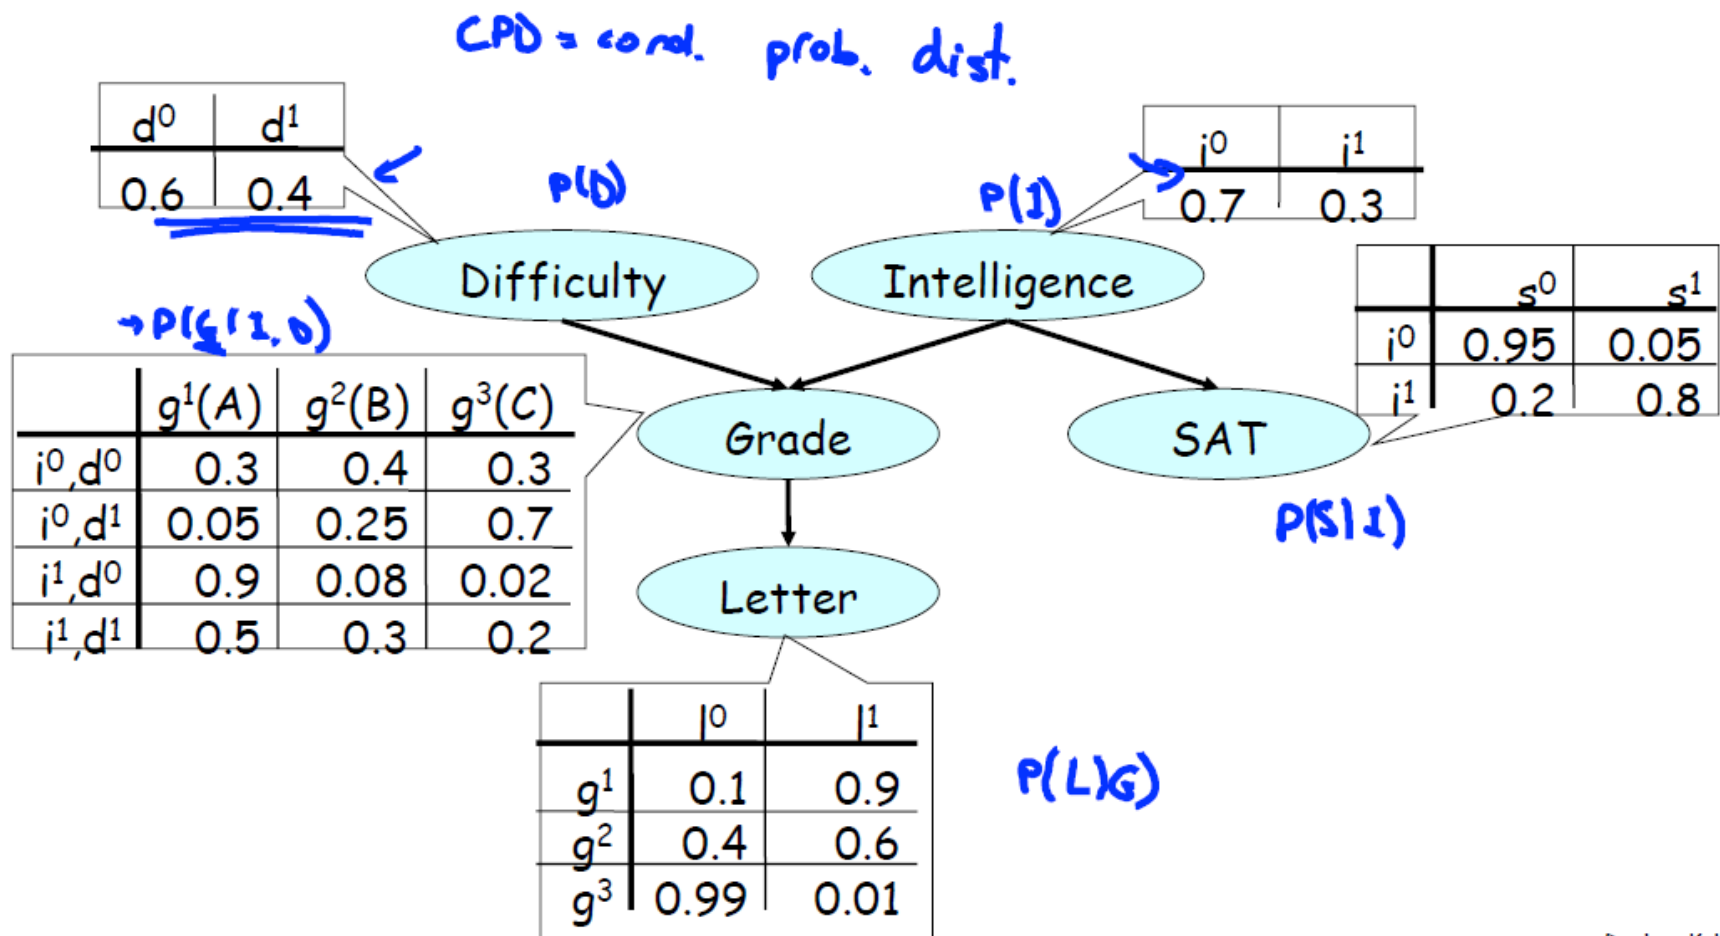
```
<Book ISBN="ISBN-0-13-713526-2" Price="85" Edition="3rd">  
  <Title>A First Course in Database Systems </Title>  
  <Authors>  
    <Auth> <FName>Jeffrey</FName> <LName>Ullman</LName> </Auth>  
    <Auth><FName>Jennifer</FName><LName>Widom</LName></Auth>  
  </Authors>  
</Book>
```

Data Model for Probability Distributions

- Probabilistic Graph Models
 - Designed for working with Complex Probability Distributions
- Example: Student Domain Variables
 - Course Difficulty (Difficulty)
 - Student grade in Course (Grade)
 - Quality of Recommendation Letter (Letter)
 - Student Intelligence (Intelligence)
 - Student SAT Score (SAT)
- Given a student's SAT score and Course Difficulty:
 - What is the probability student will receive an A.

Probabilistic Graph Models

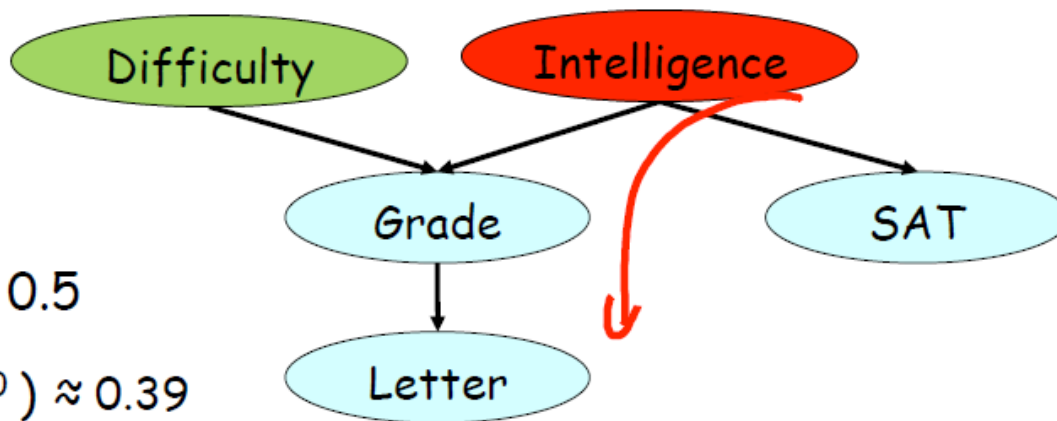
- Bayesian Networks



Probabilistic Graph Models

- Causal Inference

Causal Reasoning



$P(I^1) \approx 0.5$

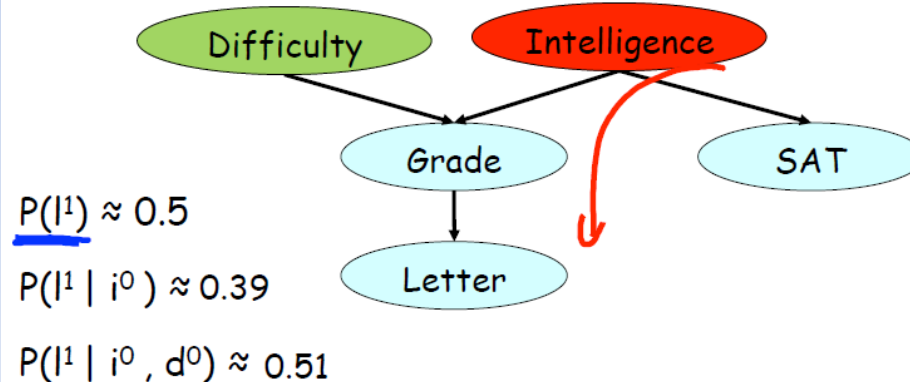
$P(I^1 | i^0) \approx 0.39$

$P(I^1 | i^0, d^0) \approx 0.51$

Probabilistic Graph Models

- Judea Pearl – Turing Award, 2011
 - For fundamental contributions to artificial intelligence through the development of a calculus for probabilistic and causal reasoning.

Causal Reasoning



Relational Data Model

- Codd 1970
- Everything is a table
- Every row in a table has the same columns
- Relationships are implicit no pointers

Database Philosophy

God made the integers;
all else is the work of man.
(Leopold Kronecker, 19th Century Mathematician)

Codd made relations;
all else is the work of man.
(Raghu Ramakrishnan, DB text book author)

Map-Reduce Model

- Unlike the Relational Data Model and other models, it is a Programming Paradigm
- Existing MapReduce and Similar Systems
 - Google MapReduce
 - Support C++, Java, Python, Sawzall, etc.
 - Based on proprietary infrastructures and some open source libraries
- Hadoop Map-Reduce
 - Open Source!
 - HDFS, Map-Reduce, Pig, Zookeeper, HBase, Hive
 - Used by Yahoo!, Facebook, Amazon and Google-IBM NSF cluster
- Dryad
 - Proprietary, based on Microsoft SQL servers

Data Model - Relations

- Structure of the Data: Relations
 - Physical Data Model
 - Conceptual Data Model
- Operations on the Data:
 - Queries
 - Modification
 - Allows Optimizations
- Constraints on the Data

Relational Data Model

- Provides limited, yet useful, operations
- Provides framework for power of languages
 - SQL

Relation Example

Person Relation

personID	firstName	lastName	gender	birthdate
100	John	Doe100	M	1960-10-10
101	John	Doe101	M	1955-01-10
102	Jane	Doe102	F	1955-10-01
103	John	Doe103	M	1970-09-10
104	Jane	Doe104	F	1977-07-07

Relational Data Model

Terminology

- Attributes
 - Columns of relation
- Schema
 - Schema of Relation= Name + {Attributes}
 - Schema of Database = {Schema of Relations}

Relational Data Model

Terminology

- Tuples:
 - Rows of relation
- Domains: Attributes have domains
- Relations are Sets of Tuples
- Instance of Relation
 - Set of tuples from relation
 - Current Instance = Set of tuples in relation NOW

Relation Example

Person Relation

personID	firstName	lastName	gender	birthdate
100	John	Doe100	M	1960-10-10
101	John	Doe101	M	1955-01-10
102	Jane	Doe102	F	1955-10-01
103	John	Doe103	M	1970-09-10
104	Jane	Doe104	F	1977-07-07

- Attributes?
 - PersonID, FirstName, LastName, Gender, Birthdate

Relation Example

Person Relation

personID	firstName	lastName	gender	birthdate
100	John	Doe100	M	1960-10-10
101	John	Doe101	M	1955-01-10
102	Jane	Doe102	F	1955-10-01
103	John	Doe103	M	1970-09-10
104	Jane	Doe104	F	1977-07-07

- Tuples?

(100, John, Doe100, M, 1960-10-10)

(101, John, Doe101, M, 1955-01-10)

(102, John, Doe102, F, 1955-10-01)

(103, John, Doe103, M, 1970-09-10)

(104, John, Doe104, F, 1977-07-07)

Relation Example

Person Relation

personID	firstName	lastName	gender	birthdate
100	John	Doe100	M	1960-10-10
101	John	Doe101	M	1955-01-10
102	Jane	Doe102	F	1955-10-01
103	John	Doe103	M	1970-09-10
104	Jane	Doe104	F	1977-07-07

- Components of the first tuple?

100 -> PersonID

John -> FirstName

Doe100 -> LastName

M -> Gender

1960-10-10 -> Birthdate

Relation Example

Person Relation

personID	firstName	lastName	gender	birthdate
100	John	Doe100	M	1960-10-10
101	John	Doe101	M	1955-01-10
102	Jane	Doe102	F	1955-10-01
103	John	Doe103	M	1970-09-10
104	Jane	Doe104	F	1977-07-07

- The Relation Schema for Person?

Person(personID, firstName, lastName, gender, birthdate)

Relation Example

Person Relation

personID	firstName	lastName	gender	birthdate
100	John	Doe100	M	1960-10-10
101	John	Doe101	M	1955-01-10
102	Jane	Doe102	F	1955-10-01
103	John	Doe103	M	1970-09-10
104	Jane	Doe104	F	1977-07-07

- Suitable Domain for each attribute
 - personID: integer
 - firstName, lastName: string (varchar in MySQL)
 - gender: Character(1)
 - birthDate: Date

Person Relation:

Equivalent Representation

personID	lastName	firstName	birthdate	gender
100	Doe100	John	1960-10-10	M
101	Doe101	John	1955-01-10	M
102	Doe102	Jane	1955-10-01	F
103	Doe103	John	1970-09-10	M
104	Doe104	Jane	1977-07-07	F

Relational Data Model

Terminology

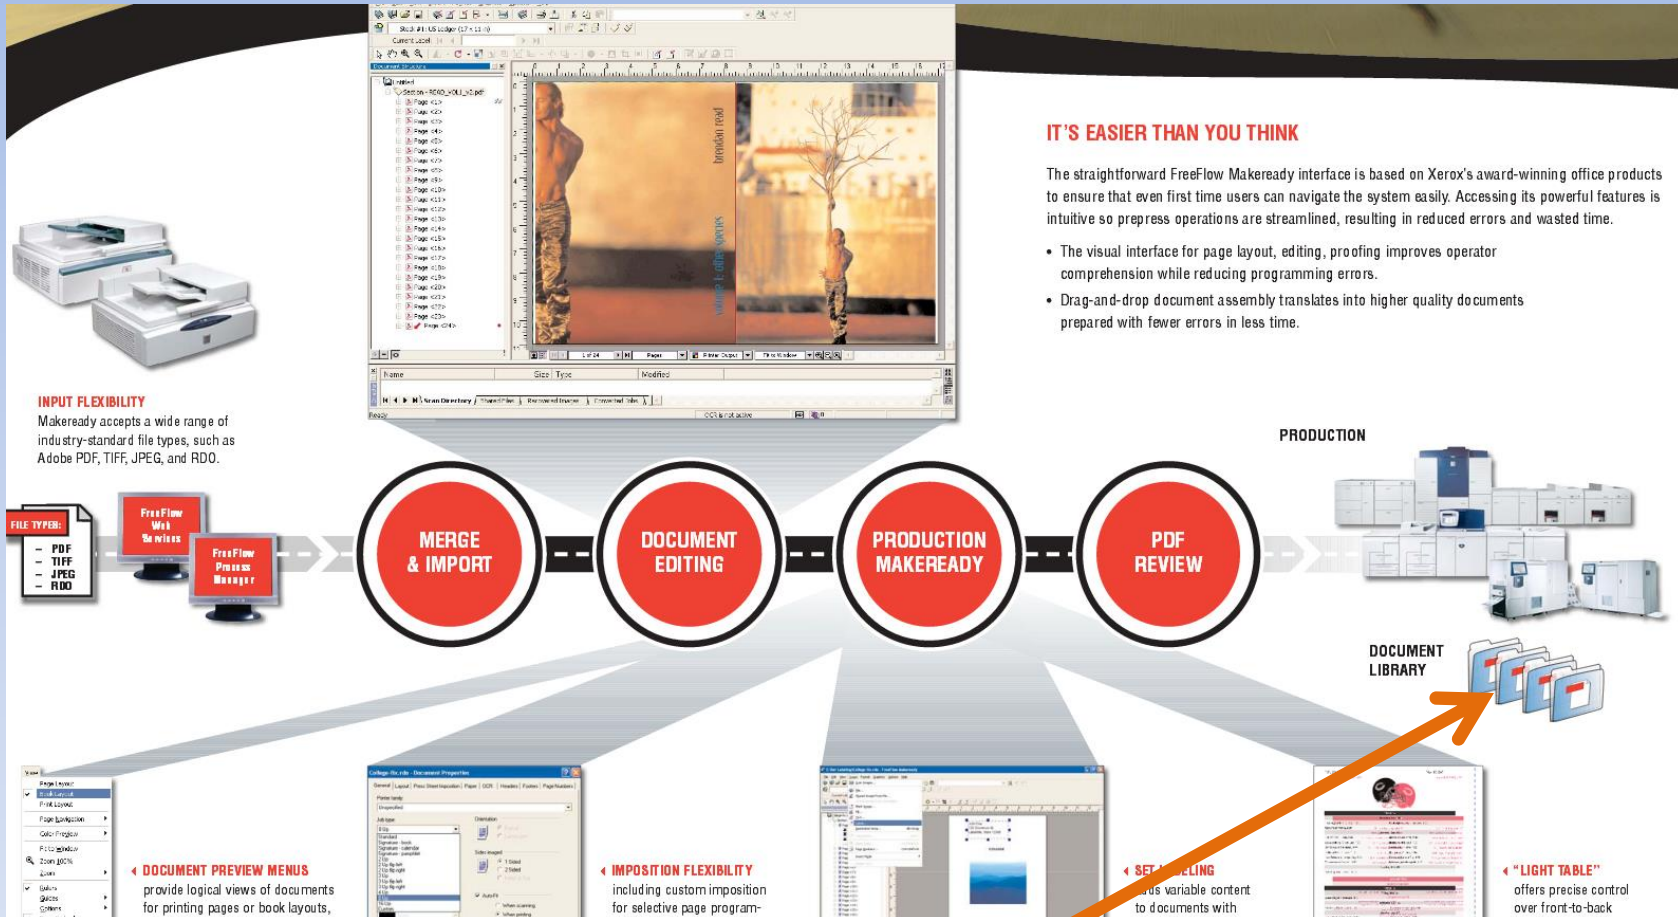
- Key Constraint
 - Attribute (or set) that are GUARANTEED unique.
- Key Examples:
 - SSN
 - Vehicle VIN
 - Student ID
 - Book ISBN

Relational Data Model: Part 2

Language: SQL (se.quel)

- Two Aspects to language
- Data Definition (ddl):
 - Declaring database schema: tables, constraints, indexes, views.
 - like declaring data/variables in programming language
- Data Manipulation (dml):
 - asking questions (querying) and modifying data.
 - Like executable code within programming language.

Database Example w/ Xerox FreeFlow



DataBase

Database Examples

Your Benefits Resources™

AON

Home

Retirement ▾

Other Benefits ▾

Knowledge Center ▾

Print

Change My Future Investments

Choose Funds

Review

Complete


You've chosen to change how your future contributions will be invested. This **won't** affect your current balance.

Tools and Information ⓘ


Savings Plan Investment Elections - All Accounts

Asset Class	Current Mix	New Mix
Fund		
Target Retirement Income Portfolios		
Target Date	20%	
TRIP Income	0%	<input type="text"/> 0%
TRIP 2010	0%	<input type="text"/> 0%
TRIP 2020	0%	<input type="text"/> 0%
TRIP 2030	10%	<input type="text"/> 0%
TRIP 2040	10%	<input type="text"/> 0%
TRIP 2050	0%	<input type="text"/> 0%
Core Funds		
GIC/Stable Value	10%	
Stable Value	10%	<input type="text"/> 0%
Bond	10%	
US Bond Index	10%	<input type="text"/> 0%
Diversified Bond	0%	<input type="text"/> 0%
Balanced	0%	
Diversified Asset Allocation	0%	<input type="text"/> 0%
Large U.S. Equity	10%	

Database Experience

YourBenefitsResources™

[Home](#)[Retirement ▼](#)[Other Benefits ▼](#)[Knowledge Center ▼](#)

 [Print](#)

Retirement Account Summary

401(k) Savings


401(k) Savings

Account Summary

As of 01-14-2014

This includes Aon Savings Plan.

Balances	Details
Current Balance	\$0.00
Rate of Return	Details
Your Rate of Return (Year-to-Date)	0.00%
Automatic Rebalancing	Start
Next Rebalance Date	Not Active

delivered by  [About This Site](#) | [Legal Info](#) | [Privacy Statement](#) | [Feedback](#) | [Contact Us](#) | [Log Off](#)

Chapter 2: Relational Algebra

- Data Manipulation Language Still Needed
- Enter Relational Algebra
 - In commercial systems not used directly
 - SQL (in commercial systems used directly) has Relational Algebra as it's center.
 - SQL query gets translated to Relational Algebra
 - Limited expressiveness a virtue.

Algebra

- What is Algebra?
 - Operands & Operators
 - Closure is usually needed.
- Let's take a look @ Some Algebras

Algebra - Arithmetic

- What is Algebra?
 - Operands & Operators
 - Closure is usually needed.
- Algebra of Arithmetic
 - Operands include: variables (x, y, z, \dots), and constants ($1, 4, 5, \dots$)
 - Operators include: addition, subtractions, division, ...

Algebra - Arithmetic

- What is Algebra?
 - Operands & Operators
 - Closure is usually needed.
- Algebra of Arithmetic
 - Operands include: variables (x, y, z, \dots), and constants ($1, 4, 5, \dots$)
 - Operators include: addition, subtractions, division, ...
- **Linear Algebra**
 - Operands: Matrix variables and constants
 - Operators include Dot Product, Determinant, Transpose
 - Operations returns matrices, allowing operator composition
 - Building Expressions

Relation Algebra

- Operands:
 - Relations
 - Variables representing relations
- Operators:
 - Set operations
 - Slicing Relations
 - Gluing Relations
 - Renaming Relations
- Closure is Needed

Example Relation Schemas

- R1(K, A, B, C)
- R2(K, D, E)
- R3(A, A1, A2, A3)
- R4(B, B1, B2)
- R5(C, C1, C2, C3, C4, C5)
- w/ K a key value for R1 and R2.
- w/ A a key value for R3.
- w/ B a key value for R4.
- w/ C a key value for R5.

Current Instances for Relation Examples

R1

K	A	B	C
4	2	0	6
5	2	0	5
1	1	3	8
2	1	3	7
3	2	3	3

R2

K	D	E
4	1	6
5	1	5
1	1	8
2	1	7
3	1	3

R4

B	B1	B2
0	0	0
3	9	27

R5

C	C1	C2	C3	C4	C5
4	2	0	6	1	6
5	2	0	5	1	5
1	1	3	8	1	8
2	1	3	7	1	7
3	2	3	3	1	3

Set Operators

Union/Intersection/Difference

- $X \cap Y$
- $X \cup Y$
- $Y - X$
 - Schemas must be identical

Operators

Union

- $X \cup Y$

K	D	E
1	1	8
2	1	7

X

K	D	E
4	1	6
5	1	5

Y

K	D	E
4	1	6
5	1	5
1	1	8
2	1	7

$X \cup Y$

Operators

Intersection

- $X \cap Y$

K	D	E
4	1	6
5	1	5
1	1	8
2	1	7

X

K	D	E
1	1	8
2	1	7
3	1	3

Y

K	D	E
1	1	8
2	1	7

$X \cap Y$

Operators

Difference

- $X - Y$:
- Set of elements in X but NOT IN Y
- Element of Y ALSO IN X are removed

K	D	E
4	1	6
5	1	5
1	1	8
2	1	7

X

K	D	E
1	1	8
2	1	7
3	1	3

Y

K	D	E
4	1	6
5	1	5

$X - Y$

Combining Operators

- Since each operation returns a Relation (closure) it can feed other operations.
- Can be viewed as an Expression Tree

Operators

Intersection as Difference

- $X \cap Y: X - (X - Y)$

K	D	E
4	1	6
5	1	5
1	1	8
2	1	7

X

K	D	E
4	1	6
5	1	5

K	D	E
1	1	8
2	1	7

K	D	E
1	1	8
2	1	7
3	1	3

Y

$X - Y$

$X - (X - Y)$

$X \cap Y$

Operators

Selection

- $Y := \sigma_C(X)$
 - Select a set of rows of a relation
 - Based on conditional expression C
 - Operands in C are either attributes of relation X or constants.
 - Y includes only tuples that make C true.

Operators Selection

- $Y := \sigma_{K < 3}(X)$

X

K	A	B	C
4	2	0	6
5	2	0	5
1	1	3	8
2	1	3	7
3	2	3	3

$\sigma_{K < 3}(X)$

K	A	B	C
1	1	3	8
2	1	3	7

Operators

Projection

- $Y := \Pi_L(X)$
 - Select a set of attributes/columns of relation

Operators

Projection

- $Y := \pi_{C,C2}(X)$

C	C1	C2	C3	C4	C5
4	2	0	6	1	6
5	2	0	5	1	5
1	1	3	8	1	8
2	1	3	7	1	7
3	2	3	3	1	3

X

C	C2
4	0
5	0
1	3
2	3
3	3

$\pi_{C,C2}(X)$

Product

- $Z := X X Y$
 - Each rows of X attached to Each Possible row of Y

Product

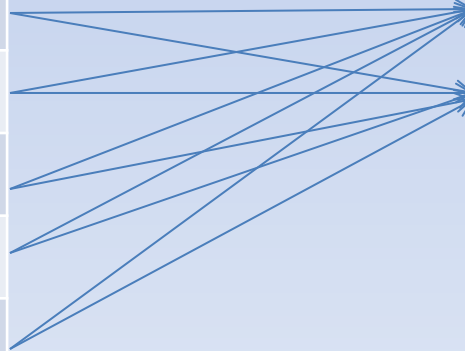
- $Z := R1 \times R4$

R1

K	A	B	C
4	2	0	6
5	2	0	5
1	1	3	8
2	1	3	7
3	2	3	3

R4

B	B1	B2
0	0	0
3	9	27



Product

R1

K	A	B	C
4	2	0	6
5	2	0	5
1	1	3	8
2	1	3	7
3	2	3	3

B	B1	B2
0	0	0
3	9	27

R4

R1 X R4

K	A	B	C	B	B1	B2
4	2	0	6	0	0	0
4	2	0	6	3	9	27
5	2	0	5	0	0	0
5	2	0	5	3	9	27
1	1	3	8	0	0	0
1	1	3	8	3	9	27
2	1	3	7	0	0	0
2	1	3	7	3	9	27
3	2	3	3	0	0	0
3	2	3	3	3	9	27

Product

$Y3 := Y1 \times Y2$

Y1(

A,	B)
1	2
3	4

Y2(

B,	C)
5	6
7	8
9	10

Y3(

A,	Y1.B,	Y2.B,	C)
1	2	5	6
1	2	7	8
1	2	9	10
3	4	5	6
3	4	7	8
3	4	9	10

Natural Join

- Usually want to join tuples that in some way *match*.
- Natural Join requires matching attributes have matching values.
- $R3 := R1 \bowtie R2$.
 - Take Product: $R1 \times R2$
 - Take Result: σ_C
 - C is same named attributes are equal
 - Remove redundant attributes
- Dangling Tuples are tuples from one relation that have no match in the other tuple.

Product – Changed Instances

$$Y3 := Y1 \times Y2$$

Y1(

A,	B)
1	2
3	4

Y2(

B,	C)
2	11
5	9
7	2

Y3(

A,	Y1.B,	Y2.B,	C)
1	2	2	11
1	2	5	9
1	2	7	2
3	4	2	11
3	4	5	9
3	4	7	2

Natural Join

$$Y3 := Y1 \bowtie Y2$$

Y1(

A,	B)
1	2
3	4

Y3(

A,	B,	C)
1	2	11

Y2(

B,	C)
2	11
5	9
7	2

Natural Join:

Product & Selection & Projection

$Y3 := Y1 \bowtie Y2$

Y1(

A,	B)
1	2
3	4

Y2(

B,	C)
2	11
5	9
7	2

Y3(

A,	Y1.B,	Y2.B,	C)
1	2	2	11
1	2	5	9
1	2	7	2
3	4	2	11
3	4	5	9
3	4	7	2

STILL NEED:

1. Take Result where matching attributes are equal: $Y4 = \sigma_{Y1.b=Y2.b} (Y3)$

2. Remove redundant attributes: $\pi_{A,Y1.B,C} (Y4)$

Natural Join:

Dangling Tuples

$Y3 := Y1 \bowtie Y2$

Y1(

A,	B)
1	2
3	4

Y2(

B,	C)
2	11
5	9
7	2

Y3(

A,	Y1.B,	Y2.B,	C)
1	2	2	11
1	2	5	9
1	2	7	2
3	4	2	11
3	4	5	9
3	4	7	2

STILL NEED:

1. Take Result where matching attributes are equal: $Y4 = \sigma_{Y1.b=Y2.b} (Y3)$

2. Remove redundant attributes: $\pi_{A,Y1.B,C} (Y4)$

Theta-Join

- $R3 := R1 \bowtie_C R2$
 - Take Product: $R1 \times R2$
 - Take Result: σ_C
 - C is boolean condition

Theta Join:

Product & Selection

$$Y3 := \sigma_{A < C} (Y1 \times Y2)$$

Y1(

A,	B)
1	2
3	4

Y2(

B,	C)
2	11
5	9
7	2

Y3(

A,	Y1.B,	Y2.B,	C)
1	2	2	11
1	2	5	9
1	2	7	2
3	4	2	11
3	4	5	9
3	4	7	2

Theta Join:

Product & Selection

$$Y3 := \sigma_{A+B < C} (Y1 \times Y2)$$

Y1(

A,	B)
1	2
3	4

Y2(

B,	C)
2	11
5	9
7	2

Y3(

A,	Y1.B,	Y2.B,	C)
1	2	2	11
1	2	5	9
1	2	7	2
3	4	2	11
3	4	5	9
3	4	7	2

Renaming

- $R1 := \rho_{R1(A1, \dots, An)}(R2)$
 - makes R1 be a relation with attributes $A1, \dots, An$ and the same tuples as R2.

Renaming

$Y3 := \rho_{R1(A,B1,B2,C)} (Y1 \times Y2)$

Y1(

A,	B)
1	2
3	4

Y2(

B,	C)
2	11
5	9
7	2

R1(

A,	B1,	B2,	C)
1	2	2	11
1	2	5	9
1	2	7	2
3	4	2	11
3	4	5	9
3	4	7	2

Renaming & Set Operations

$$Y3 := \rho_{X,Y}(Y1) \cup \rho_{X,Y}(Y2)$$

Y1(

A,	B)
1	2
3	4

Y2(

B,	C)
2	11
5	9
7	2

Y3(

X,	Y)
1	2
3	4
2	11
5	9
7	2

Precedence

- Precedence of relational operators:
 1. $[\sigma, \pi, \rho]$ (highest).
 2. $[X, \bowtie]$.
 3. \cap .
 4. $[\cup, -]$

Relational Algebra for Constraints

- Relation $R \neq \emptyset$ OR $R = \emptyset$
- Key Constraints
- Foreign Key Constraints
- Domain Value Constraints

Example 1: Key Constraint

- $PC1 = PC2 = PC$

$$Y := \sigma_{PC1.model=PC2.model \text{ AND } PC maker \neq PC2 maker}(PC1 \times PC2)$$

- Key Constraint: $Y = \emptyset$

Example 2

Referential Integrity Constraint

$Y1 := \pi_{\text{model}}(\text{Product})$

$Y2 := \pi_{\text{model}}(\text{PC})$

- Referential Integrity Constraint: $Y2 \subseteq Y1$

Example:

Exercise – 2.4.1

- Product(maker, model, type)
- PC(model, speed, ram, hd, price)
- Laptop(model, speed, ram, hd, screen, price)
- Printer(model, color, type, price)

a) What PC models have a speed of at least 3.00?

<i>maker</i>	<i>model</i>	<i>type</i>
A	1001	pc
A	1002	pc
A	1003	pc
A	2004	laptop
A	2005	laptop
A	2006	laptop
B	1004	pc
B	1005	pc
B	1006	pc
B	2007	laptop
C	1007	pc
D	1008	pc
D	1009	pc
D	1010	pc
D	3004	printer
D	3005	printer
E	1011	pc
E	1012	pc
E	1013	pc
E	2001	laptop
E	2002	laptop
E	2003	laptop
E	3001	printer
E	3002	printer
E	3003	printer
F	2008	laptop
F	2009	laptop
G	2010	laptop
H	3006	printer
H	3007	printer

<i>model</i>	<i>speed</i>	<i>ram</i>	<i>hd</i>	<i>price</i>
1001	2.66	1024	250	2114
1002	2.10	512	250	995
1003	1.42	512	80	478
1004	2.80	1024	250	649
1005	3.20	512	250	630
1006	3.20	1024	320	1049
1007	2.20	1024	200	510
1008	2.20	2048	250	770
1009	2.00	1024	250	650
1010	2.80	2048	300	770
1011	1.86	2048	160	959
1012	2.80	1024	160	649
1013	3.06	512	80	529

(a) Sample data for relation PC

<i>model</i>	<i>speed</i>	<i>ram</i>	<i>hd</i>	<i>screen</i>	<i>price</i>
2001	2.00	2048	240	20.1	3673
2002	1.73	1024	80	17.0	949
2003	1.80	512	60	15.4	549
2004	2.00	512	60	13.3	1150
2005	2.16	1024	120	17.0	2500
2006	2.00	2048	80	15.4	1700
2007	1.83	1024	120	13.3	1429
2008	1.60	1024	100	15.4	900
2009	1.60	512	80	14.1	680
2010	2.00	2048	160	15.4	2300

(b) Sample data for relation Laptop

<i>model</i>	<i>color</i>	<i>type</i>	<i>price</i>
3001	true	ink-jet	99
3002	false	laser	239
3003	true	laser	899
3004	true	ink-jet	120
3005	false	laser	120
3006	true	ink-jet	100
3007	true	laser	200

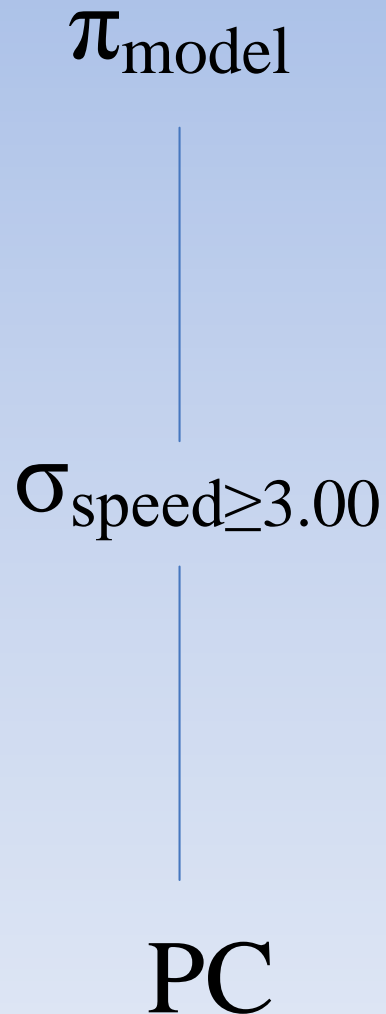
(c) Sample data for relation Printer

a) What PC models have a speed of at least 3.00?

- $R1 := \sigma_{\text{speed} \geq 3.00}(\text{PC})$
- $R2 := \pi_{\text{model}}(R1)$

model
1005
1006
1013

a) What PC models have a speed of at least 3.00?



Example:

Exercise – 2.4.1

- Product(maker, model, type)
- PC(model, speed, ram, hd, price)
- Laptop(model, speed, ram, hd, screen, price)
- Printer(model, color, type, price)

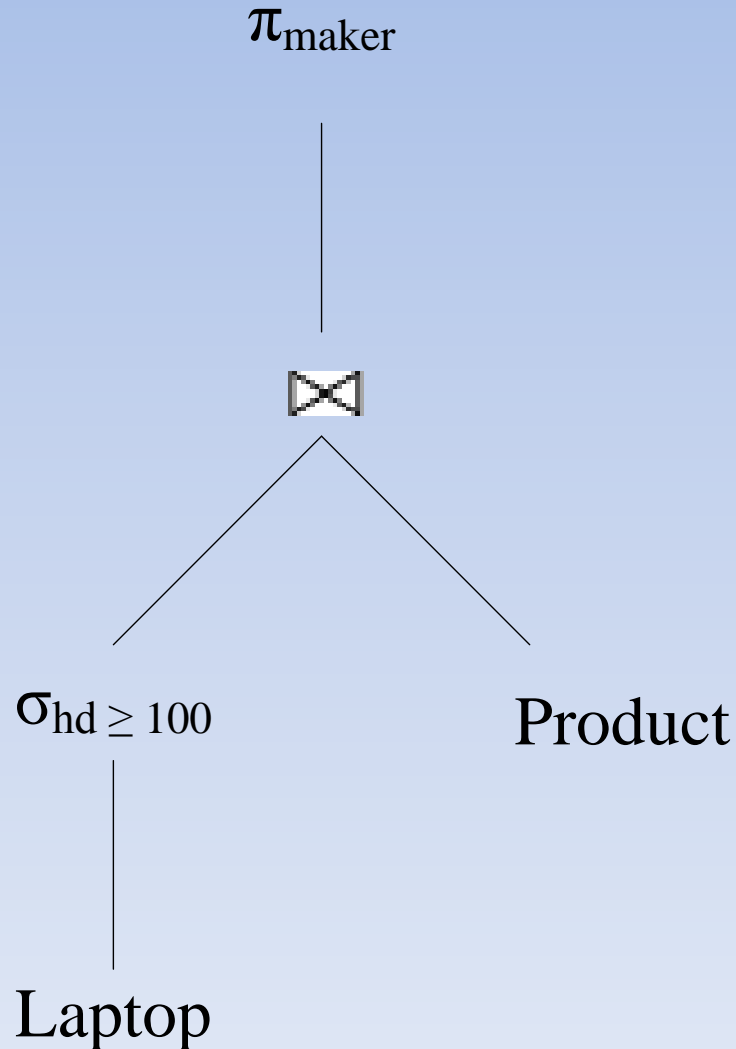
b) Which manufacturers make laptops with a hard disk of at least 100gb

b) Which manufacturers make laptops with a hard disk of at least 100gb

- $R1 := \sigma_{hd \geq 100} (\text{Laptop})$
- $R2 := \text{Product } (R1)$
- $R3 := \pi_{\text{maker}} (R2)$

maker
E
A
B
F
G

b) Which manufacturers make laptops with a hard disk of at least 100gb



Example:

Exercise – 2.4.1

- Product(maker, model, type)
 - PC(model, speed, ram, hd, price)
 - Laptop(model, speed, ram, hd, screen, price)
 - Printer(model, color, type, price)
- c) Find all model number and price of all products (of any type) made by manufacturer B

c) Find all model number and price of all products (of any type) made by manufacturer B

- $R1 := \sigma_{\text{maker}=\text{B}} (\text{Product} \bowtie \text{PC})$
- $R2 := \sigma_{\text{maker}=\text{B}} (\text{Product} \bowtie \text{Laptop})$
- $R3 := \sigma_{\text{maker}=\text{B}} (\text{Product} \bowtie \text{Printer})$
- $R4 := \pi_{\text{model}, \text{price}} (R1)$
- $R5 := \pi_{\text{model}, \text{price}} (R2)$
- $R6 := \pi_{\text{model}, \text{price}} (R3)$
- $R7 := R4 \cup R5 \cup R6$

model	price
1004	649
1005	630
1006	1049
2007	1429

In-Class 2b

Ullman & Widom pg. 52

- Product(maker, model, type)
- PC(model, speed, ram, hd, price)
- Laptop(model, speed, ram, hd, screen, price)
- Printer(model, color, type, price)

- **Write the Relational Algebra to:**

Ex: 2.4.1.d) Find the model numbers of all color laser printers

Ex: 2.4.1.e) Find those manufacturers (maker) that sell laptops, but not PC's.

maker	model	type
A	1001	pc
A	1002	pc
A	1003	pc
A	2004	laptop
A	2005	laptop
A	2006	laptop
B	1004	pc
B	1005	pc
B	1006	pc
B	2007	laptop
C	1007	pc
D	1008	pc
D	1009	pc
D	1010	pc
D	3004	printer
D	3005	printer
E	1011	pc
E	1012	pc
E	1013	pc
E	2001	laptop
E	2002	laptop
E	2003	laptop
E	3001	printer
E	3002	printer
E	3003	printer
F	2008	laptop
F	2009	laptop
G	2010	laptop
H	3006	printer
H	3007	printer

model	speed	ram	hd	price
1001	2.66	1024	250	2114
1002	2.10	512	250	995
1003	1.42	512	80	478
1004	2.80	1024	250	649
1005	3.20	512	250	630
1006	3.20	1024	320	1049
1007	2.20	1024	200	510
1008	2.20	2048	250	770
1009	2.00	1024	250	650
1010	2.80	2048	300	770
1011	1.86	2048	160	959
1012	2.80	1024	160	649
1013	3.06	512	80	529

(a) Sample data for relation PC

model	speed	ram	hd	screen	price
2001	2.00	2048	240	20.1	3673
2002	1.73	1024	80	17.0	949
2003	1.80	512	60	15.4	549
2004	2.00	512	60	13.3	1150
2005	2.16	1024	120	17.0	2500
2006	2.00	2048	80	15.4	1700
2007	1.83	1024	120	13.3	1429
2008	1.60	1024	100	15.4	900
2009	1.60	512	80	14.1	680
2010	2.00	2048	160	15.4	2300

(b) Sample data for relation Laptop

model	color	type	price
3001	true	ink-jet	99
3002	false	laser	239
3003	true	laser	899
3004	true	ink-jet	120
3005	false	laser	120
3006	true	ink-jet	100
3007	true	laser	200

(c) Sample data for relation Printer

- Write the Relational Algebra to:

- Ex: 2.4.1.d) Find the model numbers of all color laser printers
- Ex: 2.4.1.e) Find those manufacturers (maker) that sell laptops, but not PC's.

Answers Follow

In-Class 2b

- Product(maker, model, type)
- PC(model, speed, ram, hd, price)
- Laptop(model, speed, ram, hd, screen, price)
- Printer(model, color, type, price)
- **Write the Relational Algebra to:**
 - Ex: 2.4.1.d) Find the model numbers of all color laser printers
 - Ex: 2.4.1.e) Find those manufacturers (maker) that sell laptops, but not PC's.

In-Class 2b

- Product(maker, model, type)
- PC(model, speed, ram, hd, price)
- Laptop(model, speed, ram, hd, screen, price)
- Printer(model, color, type, price)
- **Write the Relational Algebra to:**

Ex: 2.4.1.d) Find the model numbers of all color laser printers

$R1 := \sigma_{\text{color} = \text{true} \text{ AND } \text{type} = \text{laser}} (\text{Printer})$

$R2 := \pi_{\text{model}} (R1)$

CORRECT

model
3003
3007

In-Class 2b

- Product(maker, model, type)
- PC(model, speed, ram, hd, price)
- Laptop(model, speed, ram, hd, screen, price)
- Printer(model, color, type, price)
- **Write the Relational Algebra to:**

Ex: 2.4.1.d) Find the model numbers of all color laser printers

$\pi_{\text{model}} (\sigma_{\text{color} = \text{true AND type} = \text{laser}} (\text{Printer}))$

CORRECT

model
3003
3007

In-Class 2b

- Product(maker, model, type)
- PC(model, speed, ram, hd, price)
- Laptop(model, speed, ram, hd, screen, price)
- Printer(model, color, type, price)
- **Write the Relational Algebra to:**

Ex: 2.4.1.d) Find the model numbers of all color laser printers

$R1 := \sigma_{\text{color} = \text{true} \text{ AND } \text{type} = \text{laser}}(\text{Printer})$

$R2 := \pi_{\text{model}}(R1)$

Partially Correct

model
3003
3007

In-Class 2b

- Product(maker, model, type)
- PC(model, speed, ram, hd, price)
- Laptop(model, speed, ram, hd, screen, price)
- Printer(model, color, type, price)
- **Write the Relational Algebra to:**

Ex: 2.4.1.d) Find the model numbers of all color laser printers

$R1 := \sigma_{\text{color}=\text{true}}(\text{Printer})$

$R2 := \sigma_{\text{type}=\text{laser}}(\text{Printer})$

$R3 := R1 \bowtie R2$

~~$R3 := R1 \cap R2$~~

~~$R4 := \pi_{\text{model}}(R3)$~~

Partially Correct

model
3003
3007

In-Class 2b

- Product(maker, model, type)
- PC(model, speed, ram, hd, price)
- Laptop(model, speed, ram, hd, screen, price)
- Printer(model, color, type, price)
- **Write the Relational Algebra to:**

Ex: 2.4.1.d) Find the model numbers of all color laser printers

$R1 := \sigma_{\text{color}=\text{true}}(\text{Printer})$

Partially Correct

model
3003
3007

In-Class 2b

- Product(maker, model, type)
- PC(model, speed, ram, hd, price)
- Laptop(model, speed, ram, hd, screen, price)
- Printer(model, color, type, price)
- **Write the Relational Algebra to:**

Ex: 2.4.1.d) Find the model numbers of all color laser printers R1

$R1 := \pi_{\text{model,color}}(\text{Printer})$

$R2 := \sigma_{\text{model}}(R1)$

Incorrect

model	Color
3002	False
3003	True
3007	True

In-Class 2b

- Product(maker, model, type)
- PC(model, speed, ram, hd, price)
- Laptop(model, speed, ram, hd, screen, price)
- Printer(model, color, type, price)
- **Write the Relational Algebra to:**

Ex: 2.4.1.e) Find those manufacturers (maker) that sell laptops, but not PC's.

$R1 := \sigma_{\text{type=laptop}}(\text{Product})$

$R2 := \sigma_{\text{type=PC}}(\text{Product})$

$R3 := \pi_{\text{maker}}(R1)$

$R4 := \pi_{\text{maker}}(R2)$

$R5 := R3 - R4$

CORRECT

maker
F
G

In-Class 2b

- Product(maker, model, type)
- PC(model, speed, ram, hd, price)
- Laptop(model, speed, ram, hd, screen, price)
- Printer(model, color, type, price)
- **Write the Relational Algebra to:**

Ex: 2.4.1.e) Find those manufacturers (maker) that sell laptops, but not PC's.

$R1 := \sigma_{\text{type}=\text{laptop}}(\text{Product})$

$R2 := \sigma_{\text{type}=\text{PC}}(\text{Product})$

$\pi_{\text{maker}}(\text{R1}-\text{R2})$

Correct: $\pi_{\text{maker}}(R1)-\pi_{\text{maker}}(R2)$

R1

maker	model	type
A	2004	laptop
A	2006	laptop
F	2008	laptop

R2

maker	model	type
A	1002	PC
A	1003	PC

In-Class 2b

- Product(maker, model, type)
- PC(model, speed, ram, hd, price)
- Laptop(model, speed, ram, hd, screen, price)
- Printer(model, color, type, price)
- **Write the Relational Algebra to:**

Ex: 2.4.1.e) Find those manufacturers (maker) that sell laptops, but not PC's.

~~$R1 := \sigma_{\text{type=laptop}}(\text{Product})$~~

~~$R2 := \sigma_{\text{type=PC}}(\text{Product})$~~

$R3 := \sigma_{\text{maker}}(\text{laptop})$

$R4 := \sigma_{\text{maker}}(\text{PC})$

$R5 := R3 - R4$

Partially Correct

maker	
F	
G	

DB4

Courseware

Course Info

Discussion

Wiki

Progress

Readings

Software Guide

Extra Problems

▸ Getting Started

▼ Relational Algebra

Select, Project, Join

**Set Operators, Renaming,
Notation**

Relational Algebra Quiz
Quiz



**Relational Algebra
Exercises**
Exercise



▸ Course Completion



In this assignment you are to write relational algebra queries over a small database, executed using our RA Workbench. Behind the scenes, the RA workbench translates relational algebra expressions into SQL queries over the database stored in SQLite. Since relational algebra symbols aren't readily available on most keyboards, RA uses a special syntax described in our [RA Relational Algebra Syntax](#) guide.

We've created a small sample database to use for this assignment. It contains four relations:

```
Person(name, age, gender)      // name is a key
Frequents(name, pizzeria)      // [name,pizzeria] is a key
Eats(name, pizza)              // [name,pizza] is a key
Serves(pizzeria, pizza, price) // [pizzeria,pizza] is a key
```


RA: A Relational Algebra Interpreter

<http://www.cs.duke.edu/~junyang/ra/>

Introduction

RA is a simple relational algebra interpreter written in Java. It is built on top of an SQL-based relational database system. It implements relational algebra queries by translating them into SQL queries and executing them on the underlying database system through JDBC. RA is packaged with SQLiteJDBC, so you can use RA as a standalone relational-algebra database system. Alternatively, you can use RA as a relational-algebra frontend to other database systems.

RA Operators

- $\backslash\text{select_}\{cond\}$
- $\backslash\text{project_}\{attr_list\}$
- $\backslash\text{join_}\{cond\}$
- $\backslash\text{join}$
- $\backslash\text{cross}$
- $\backslash\text{union}$, $\backslash\text{diff}$, and $\backslash\text{intersect}$ are the relational union, difference, and intersect operators.
- $\backslash\text{rename_}\{new_attr_name_list\}$ is the relational rename operator,

RA Operators

- `\select_{cond}`
 - is the relational selection operator.
 - For example, to select Drinker tuples with name Amy or Ben, we can write
 - `\select_{name = 'Amy' or name = 'Ben'} Drinker;`
 - Syntax for *cond* follows SQL.
 - Note that string literals should be enclosed in **single** quotes, and you may use boolean operators and, or, and not.
 - Comparison operators `<=`, `<`, `=`, `>`, `>=`, and `<>` work on both string and numeric types.
 - For string match you can use the SQL LIKE operator;
 - `\select_{name like 'A%'} drinker;`
 - finds all drinkers whose name start with A, as % is a wildcard character that matches any number of characters.

RA Operators

- $\text{\textbackslash project_}\{attr_list\}$
 - is the relational projection operator,
 - *attr_list* is a comma-separated list of attribute names.
 - For example, to find out what beers are served by Talk of the Town (but without the price information), we can write :

$\text{\textbackslash project_}\{bar, beer\} (\text{\textbackslash select_}\{bar = 'Talk of the Town'\} \text{Serves});$

RA Operators

- `\join_{cond}`
 - is the relational theta-join operator.
 - For example, to join `Drinker(name, address)` and `Frequents(drinker, bar, times_a_week)` relations together using drinker name, we can write
 - `Drinker \join_{name = drinker} Frequents;`
 - Syntax for *cond* again follows SQL; see notes on `\select` for more details.
- `\join`
 - is the relational natural join operator.
 - For example, to join `Drinker(name, address)` and `Frequents(drinker, bar, times_a_week)` relations together using drinker name, we can write
 - `Drinker \join \rename_{name, bar, times_a_week} Frequents;`
 - Natural join will automatically equate all pairs of identically named attributes from its inputs (in this case, `name`), and output only one attribute per pair.
 - Here we use `\rename` to create two name attributes for the natural join; see notes on `\rename` below for more details.
- `\cross`
 - is the relational cross product operator.
 - For example, to compute the cross product of `Drinker` and `Frequents`, we can write `Drinker`
 - `\cross Frequents;`

RA Operators

- \setminus union, \setminus diff, and \setminus intersect are the relational union, difference, and intersect operators.
 - For a trivial example, to compute the union, difference, and intersection between Drinker and itself, we can write
 - $\text{Drinker } \setminus \text{union Drinker};$
 - which would return Drinker itself
 - $\text{Drinker } \setminus \text{diff Drinker};$
 - an empty relation
 - $\text{Drinker } \setminus \text{intersect Drinker};$
 - Drinker itself
- $\setminus \text{rename}_{\{new_attr_name_list\}}$
 - is the relational rename operator, where *new_attr_name_list* is a comma-separated list of new names, one for each attribute of the input relation.
 - For example, to rename the attributes of relation Drinker and compute the cross product of Drinker and itself, we can write
 - $\setminus \text{rename}_{\{name1, address1\}} \text{Drinker } \setminus \text{cross } \setminus \text{rename}_{\{name2, address2\}} \text{Drinker};$

Example RA Expression

Here is an example of a complex query, which returns beers liked by those drinkers who do not frequent James Joyce Pub:

```
\project_{beer} (  
  ((\project_{name}           // all drinkers  
    Drinker)  
  \diff  
  (\rename_{name}           // rename so we can diff  
    \project_{drinker}      // drinkers who frequent JJP  
    \select_{bar = 'James Joyce Pub'}  
    Frequent))  
  
  \join_{drinker = name}    /* join with Likes to find beers */  
  
  Likes
```


Example 2: RA Expression

We've created a small sample database to use for this assignment. It contains four relations:

```
Person(name, age, gender)      // name is a key
Frequents(name, pizzeria)      // [name,pizzeria] is a key
Eats(name, pizza)              // [name,pizza] is a key
Serves(pizzeria, pizza, price) // [pizzeria,pizza] is a key
```

[View the database.](#) (You can also [download the schema and data.](#))

```
1 \project_{pizza, pizzeria}(Serves)
```

 Incorrect

Q1 (1 point possible)

Find all pizzas eaten by at least one female over the age of 20.

- [View the RA Relational Algebra Syntax guide](#)
- If you generate an error, you will see the message from the underlying SQLite system -- apologies for the lack of better error messages

```
1 \project_{pizza, pizzeria}(Serves)
```

Incorrect

Your Query Result:

cheese	Chicago Pizza
cheese	Dominos
cheese	Little Caesars
cheese	New York Pizza
cheese	Pizza Hut
cheese	Straw Hat
mushroom	Dominos
mushroom	Little Caesars
pepperoni	Little Caesars
pepperoni	New York Pizza
pepperoni	Pizza Hut
pepperoni	Straw Hat
sausage	Little Caesars

Expected Query Result:

cheese
mushroom
supreme

Reset

Submit

Q9 (1 point possible)

Find all pizzerias that serve every pizza eaten by people over 30.

(This query is very challenging; extra congratulations if you get it right.)

- [View the RA Relational Algebra Syntax guide](#)
- If you generate an error, you will see the message from the underlying SQLite system -- apologies for the lack of better error messages

```
1 Enter your RA query here
```