

Chapter 5

Chapter 5

Algebraic and Logical Query Languages

We now switch our attention to databases. We start in this section with algebraic query languages, one algebraic at a time.

206 CHAPTER 5. ALGEBRAIC AND LOGICAL QUERY LANGUAGES

<i>A</i>	<i>B</i>
1	2
3	4
1	2
1	2

Figure 5.1: A bag

5.1.1 Why Bags?

Bag Semantics

- Bags (or Multisets):
 - Generalization of the notion of a set.
 - Members are allowed to appear more than once.
- Commercial DB's implement relations as Bags.
- Some relational operations are much more efficient when done on bags.

Projection w/ Bags

R1

K	A	B	C
4	2	0	6
5	2	0	5
1	1	3	8
2	1	3	7
3	2	3	3

- $\pi_{A,B}(R1)$ -- efficiency

A	B
2	0
2	0
1	3
1	3
2	3

Bag Semantics

- Bags let us calculate some types of values more naturally.
 - Find the average B value
 - Projection onto B with sets yields (0,3)
 - Average = $3/2 = 1.5$
 - Projection onto B with bags yields (0, 0, 3, 3, 3)
 - Average = $9/5 = 1.8$

A	B
2	0
2	0
1	3
1	3
2	3

Bag Semantics

Selection

- Selection – Applied to each tuple independently.

Bag Semantics

Set Operations

- Union – Tuples from each relation added to bag.
- Intersection – Tuples in output are the min of the number of times they appear in each relation.
- Difference – Tuples are removed one for one with a minimum value of 0.

Bag Semantics

Joins

- Product – Tuples treated independently
- Joins – With multiple tuples having same values, we get multiple join possibilities.

Bag Semantics

Example

- create table t1 (x char(1), y int);
- create table t2 (x char(1), z int);
- insert into t1 values ('A',2), ('B',3), ('C',4), ('B',7);
- insert into t2 values ('B', 0), ('C',1), ('B', 4);

x	y
A	2
B	3
C	4
B	7

t1

x	z
B	0
C	1
B	4

t2

Bag Semantics

Joins

- select * from t1 natural join t2;

x	y	z
B	3	0
B	3	4
C	4	1
B	7	0
B	7	4

x	y
A	2
B	3
C	4
B	7

t1

x	z
B	0
C	1
B	4

t2

Bag Semantics – Slightly Modified Joins

- select * from t1 natural join t2;

x	y	z
B	3	0
B	3	4
C	4	1
B	3	0
B	3	4

x	y
A	2
B	3
C	4
B	3

t1

x	z
B	0
C	1
B	4

t2

Relational Algebra – Advanced (Extended)

δ = eliminate duplicates from bags.

τ = sort tuples.

γ = grouping and aggregation.

Outerjoin :

avoids “**dangling tuples**” = tuples that do not join with anything.

Duplicate Elimination

- $R1 = \delta(R2)$
 - Out relation R1 with a single copy of each tuple in R2.

Duplicate Elimination

- $R1 = \delta(R2)$

R2

x	y	z
B	3	0
B	3	4
C	4	1
B	3	0
B	3	4

$R1 = \delta(R2)$

x	y	z
B	3	0
C	4	1
B	3	4

Sort Tuples

- \mathcal{T} = sort tuples
- $R1 := \mathcal{T}_L (R2)$.
 - L is a list of attributes from $R2$.
- $R1$:
 - List of tuples of $R2$
 - sorted first on the value of the first attribute on L
 - Sorted second on the second attribute of L
 -

Sort Tuples

- $R1 := \tau_{x,z}(R2).$

R2

x	y	z
B	3	0
C	4	1
B	3	4

$\tau_{x,z}(R2)$

x	y	z
B	3	0
B	3	4
C	4	1

Aggregation Operators

- Aggregation Operators:
 - are applied to a entire column of data.
 - Return a single value.
 - i.e., SUM, AVG, COUNT, MIN, and MAX.

Aggregation Operators

- Aggregation Operators:

- SUM(K) = 15
- AVG(K) = 3
- COUNT(K) = 5
- MIN(K) = 1
- MAX(K) = 5

R1

K	A	B	C
4	2	0	6
5	2	0	5
1	1	3	8
2	1	3	7
3	2	3	3

γ – Grouping

- γ_L : (lowercase gamma)
 - Grouping Operator
- L is list of:
 - Individual Attributes (used for grouping)
 - Grouping Operators
 - (i.e., COUNT(), SUM(), ...)

$\gamma_L(R)$ – Grouping

- Group tuples in R:
 - Form one group for every set of values of attributes from L in R.
 - For each aggregation operator in AGG() in L, apply AGG() to each group formed.
- Outputs:
 - One tuple for each set of values of attributes from L in R.
 - A single value for each AGG() applied to that group.

Grouping: $\gamma_{B, count(*)}(R1)$

R1

K	A	B	C
4	2	0	6
5	2	0	5
1	1	3	8
2	1	3	7
3	2	3	3

- $\gamma_{B, count(*)}(R1)$
- Group tuples in R1:
 - Form one group for every set of values of attribute 'B' in R1 (0, 3).
 - Apply Count(*) to each group formed.

Grouping: $\gamma_{B, count(*)}(R1)$

R1

K	A	B	C
4	2	0	6
5	2	0	5
1	1	3	8
2	1	3	7
3	2	3	3

- $\gamma_B(R1)$
- Group tuples in R1:
 - Form one group for every set of values of attribute 'B' in R1 (0, 3).
 - Apply Count(*) to each group formed.

Grouping: $\gamma_{B, count(*)}(R1)$

R1

B	Count(*)
0	2
3	3

- $\gamma_{B, count(*)}(R1)$
- Group tuples in R1:
 - Form one group for every set of values of attribute 'B' in R1 (0, 3).
 - Apply Count(*) to each group formed.

Outerjoins

- $R1 \bowtie R2$
- Dangling Tuple:
 - A tuple of R1 with no corresponding tuple from R2 is said to be dangling.
 - A tuple of R2 with no corresponding tuple from R1 is also said to be dangling
 - Outerjoin preserves these tuples by padding them with null.

Outerjoins

- movie_list (mid, myear, mname);
- movie_ratings(pid, mid, rating);
- movie_list ⋈ movie_ratings

mid	myear	mname
979	1979	Movie 1
1079	1979	Movie 2
393	1993	Movie 3
1293	1993	Movie 4
300	2004	Movie 5

pid	mid	Rating
200	979	2
200	393	2
304	300	4

- movie_ratings

- movie_list

Outerjoins

- movie_list ⋈ movie_ratings

mid	myear	mname
979	1979	Movie 1
1079	1979	Movie 2
393	1993	Movie 3
1293	1993	Movie 4
300	2004	Movie 5

pid	mid	Rating
200	979	2
200	393	2
304	300	4

- movie_ratings

- movie_list

mid	myear	mname	pid	mid	Rating
979	1979	Movie 1	200	979	2
393	1993	Movie 3	200	393	2
300	2004	Movie 5	304	300	4

Outerjoins

- movie_list OUTERJOIN movie_ratings

mid	myear	mname
979	1979	Movie 1
1079	1979	Movie 2
393	1993	Movie 3
1293	1993	Movie 4
300	2004	Movie 5

pid	mid	Rating
200	979	2
200	393	2
304	300	4

- movie_ratings

- movie_list

mid	myear	mname	pid	mid	Rating
979	1979	Movie 1	200	979	2
1079	1979	Movie 2	null	null	null
393	1993	Movie 3	200	393	2
1293	1993	Movie 4	null	null	null
300	2004	Movie 5	304	300	4

SQL - Advanced

- Outer Joins
- Aggregations
- Eliminating Duplicates
- Grouping
- Having Clause
- Database Modifications

Joins - Revisited

- Join On (Theta Join)

SELECT *

FROM relation1 JOIN relation2 ON <condition1>

WHERE <condition2>;

Joins - Revisited

- Join Using

SELECT *

FROM relation1 JOIN relation2 USING (att1, ...)

WHERE <condition>;

Outerjoins w/ SQL

- `SELECT * from R1 NATURAL OUTER JOIN R2;`
- Dangling Tuple:
 - A tuple of R1 with no corresponding tuple from R2 is said to be dangling.
 - A tuple of R2 with no corresponding tuple from R1 is also said to be dangling
 - Outer Join preserves these tuples by padding them with null.

Outerjoins

- Full outer join not available in MySQL - emulated
- `SELECT * from R1 NATURAL LEFT JOIN R2 UNION
SELECT * FROM R1 NATURAL RIGHT JOIN R2 ;`
- Dangling Tuple:
 - A tuple of R1 with no corresponding tuple from R2 is said to be dangling. Left Join preserves these.
 - A tuple of R2 with no corresponding tuple from R1 is also said to be dangling. Right Join preserves these.
 - Tuples preserved by padding them with null.
 - NOTE: Second Query needs extra where clause.
 - (Reader Exercise)

Outerjoins

- movie_list (mid, myear, mname);
- movie_ratings(pid, mid, rating);
- Select * from
movie_list natural join movie_ratings

mid	myear	mname
979	1979	Movie 1
1079	1979	Movie 2
393	1993	Movie 3
1293	1993	Movie 4
300	2004	Movie 5

pid	mid	Rating
200	979	2
200	393	2
304	300	4

- movie_ratings

- movie_list

Outerjoins

- Select * from movie_list natural join movie_ratings;

mid	myear	mname
979	1979	Movie 1
1079	1979	Movie 2
393	1993	Movie 3
1293	1993	Movie 4
300	2004	Movie 5

pid	mid	Rating
200	979	2
200	393	2
304	300	4

- movie_ratings

- movie_list

mid	myear	mname	pid	mid	Rating
979	1979	Movie 1	200	979	2
393	1993	Movie 3	200	393	2
300	2004	Movie 5	304	300	4

Outerjoins

- Select * from movie_list LEFT JOIN movie_ratings;

mid	myear	mname
979	1979	Movie 1
1079	1979	Movie 2
393	1993	Movie 3
1293	1993	Movie 4
300	2004	Movie 5

pid	mid	Rating
200	979	2
200	393	2
304	300	4

- movie_ratings

- movie_list

mid	myear	mname	pid	mid	Rating
979	1979	Movie 1	200	979	2
1079	1979	Movie 2	null	null	null
393	1993	Movie 3	200	393	2
1293	1993	Movie 4	null	null	null
300	2004	Movie 5	304	300	4

Relational Algebra – Extended Now In SQL

δ = eliminate duplicates from bags.

τ = sort tuples.

γ = grouping and aggregation.

Outerjoin : avoids “dangling tuples” = tuples that do not join with anything.

Bag Semantics

Example

- create table t1 (x char(1), y int);
- create table t2 (x char(1), z int);
- insert into t1 values ('A',2), ('B',3), ('C',4), ('B',3);
- insert into t2 values ('B', 0), ('C',1), ('B', 4);

x	y
A	2
B	3
C	4
B	3

t1

x	z
B	0
C	1
B	4

t2

Bag Semantics

Joins

- `select * from t1 natural join t2;`

R2

x	y	z
B	3	0
B	3	4
C	4	1
B	3	0
B	3	4

x	y
A	2
B	3
C	4
B	7

t1

x	z
B	0
C	1
B	4

t2

Duplicate Elimination: Distinct

- SELECT * FROM R2;

R2

x	y	z
B	3	0
B	3	4
C	4	1
B	3	0
B	3	4

SELECT DISTINCT * FROM R2;

$\delta_{(R2)}$

x	y	z
B	3	0
C	4	1
B	3	4

Sort Tuples

- \mathbf{T} = sort tuples
- $R1 := \mathbf{T}_L (R2).$
 - L is a list of attributes from $R2$.
- $R1$:
 - List of tuples of $R2$
 - sorted first on the value of the first attribute on L
 - Sorted second on the second attribute of L
 -

Sort Tuples: Order By

SELECT DISTINCT * FROM R2;

x	y	z
B	3	0
C	4	1
B	3	4

SELECT DISTINCT * FROM R2 ORDER BY x,z;

$\tau_{x,z}(R2)$

x	y	z
B	3	0
B	3	4
C	4	1

Aggregation Operators

- Aggregation Operators:
 - are applied to a entire column of data.
 - Return a single value.
 - i.e., SUM, AVG, COUNT, MIN, and MAX.

Aggregation Operators

- Aggregation Operators:

- SUM(K) = 15
- AVG(K) = 3
- COUNT(K) = 5
- MIN(K) = 1
- MAX(K) = 5

R1

K	A	B	C
4	2	0	6
5	2	0	5
1	1	3	8
2	1	3	7
3	2	3	3

SELECT sum(k), avg(k), count(k),
min(k), max(k)
FROM R1;

sum(k)	avg(k)	count(k)	min(k)	max(k)
15	3	5	1	5

γ – Grouping

- γ_L : (lowercase gamma)
 - Grouping Operator
- L is list of:
 - Individual Attributes (used for grouping)
 - Grouping Operators
 - (i.e., COUNT(), SUM(), ...)

$\gamma_L(R)$ – Grouping

- Group tuples in R:
 - Form one group for every set of values of attributes from L in R.
 - For each aggregation operator in AGG() in L, apply AGG() to each group formed.
- Outputs:
 - One tuple for each set of values of attributes from L in R.
 - A single value for each AGG() applied to that group.

Grouping: $\gamma_{B, count(K)}(R1)$

R1

K	A	B	C
4	2	0	6
5	2	0	5
1	1	3	8
2	1	3	7
3	2	3	3

- $\gamma_B(R1)$
- Group tuples in R1:
 - Form one group for every set of values of attribute 'B' in R1 (0, 3).
 - Apply Count(K) to each group formed.

Grouping: $\gamma_{B, count(K)}(R1)$

R1

K	A	B	C
4	2	0	6
5	2	0	5
1	1	3	8
2	1	3	7
3	2	3	3

- $\gamma_{B, count(*)}(R1)$
- Group tuples in R1:
 - Form one group for every set of values of attribute 'B' in R1 (0, 3).
 - Apply Count(K) to each group formed.

Grouping: $\gamma_{B, count(K)}(R1)$

R1

K	A	B	C
4	2	0	6
5	2	0	5
1	1	3	8
2	1	3	7
3	2	3	3

- $\gamma_{B, count(*)}(R1)$:
 - SELECT B, Count(K) FROM R1
GROUP BY B;
 - Group tuples in R1:
 - Form one group for every set of values of attribute 'B' in R1 (0, 3).
 - Apply Count(K) to each group formed.

Grouping: $\gamma_{B, count(*)}(R1)$

R1

B	Count(K)
0	2
3	3

- $\gamma_{B, count(*)}(R1)$
 - SELECT B, Count(K) FROM R1
GROUP BY B;
 - Group tuples in R1:
 - Form one group for every set of values of attribute 'B' in R1 (0, 3).
 - Apply Count(K) to each group formed.

Iris Domain



- Iris Versicolor



- Iris Setosa



- Iris Virginica

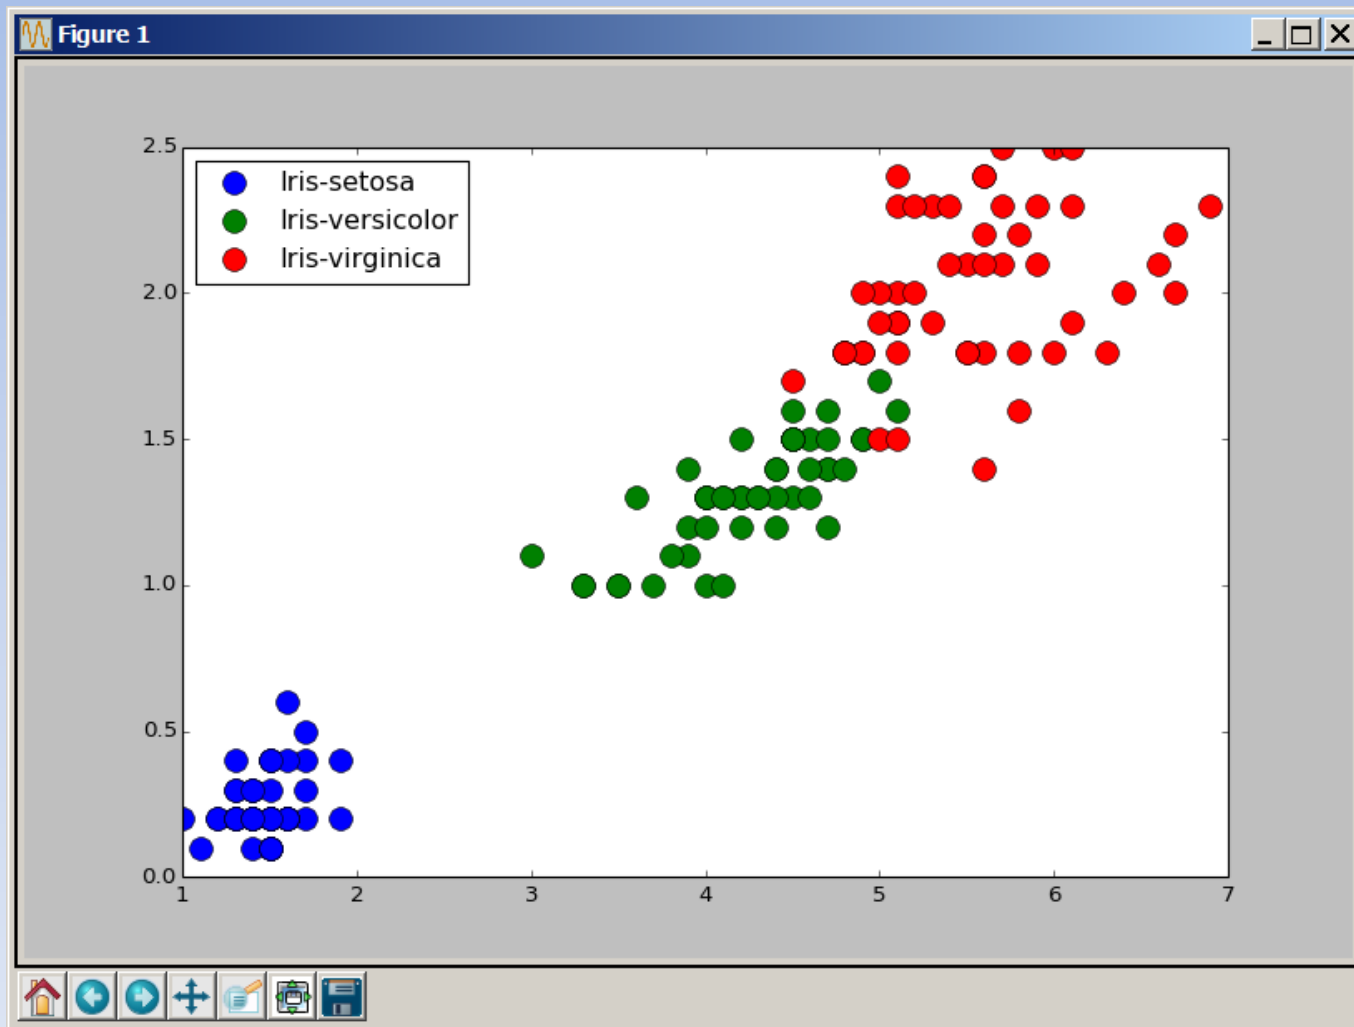
Iris_2d

petallength	petalwidth	class
1.4	0.2	Iris-setosa
1.4	0.2	Iris-setosa
1.3	0.2	Iris-setosa
1.5	0.2	Iris-setosa
1.4	0.2	Iris-setosa
3.9	1.1	Iris-versicolor
4.8	1.8	Iris-versicolor
4	1.3	Iris-versicolor
4.9	1.5	Iris-versicolor
4.7	1.2	Iris-versicolor
5.2	2.3	Iris-virginica
5	1.9	Iris-virginica
5.2	2	Iris-virginica
5.4	2.3	Iris-virginica
5.1	1.8	Iris-virginica

Iris_2d Domain

```
def test2():  
    db = "iris2d"  
    table = "iris_2d"  
    x_item = 'petallength'  
    y_item = 'petalwidth'  
    plot(db, table, x_item, y_item, 'Iris-setosa')  
    plot(db, table, x_item, y_item, 'Iris-versicolor')  
    plot(db, table, x_item, y_item, 'Iris-virginica')  
    matplotlib.pyplot.legend(loc='upper left', numpoints = 1)  
    matplotlib.pyplot.show()
```

Iris_2d Domain



Iris_2d Domain

- iris_2d (petallength float, petalwidth float, class varchar(30));
- FIND: Average petal length and width for Iris Setosa, Iris Versicolor, Iris Virginica
 - SELECT class, avg(petallength), avg(petalwidth)
FROM iris_2d
GROUP BY class
ORDER BY avg(petallength), avg(petalwidth);

class	avg(petallength)	avg(petalwidth)
Iris-setosa	1.463999996	0.244000005
Iris-versicolor	4.259999981	1.325999992
Iris-virginica	5.551999989	2.025999978

Having Clause

- HAVING <condition>
 - Can follow a GROUP BY clause
 - Condition is applied to each group
 - Groups not satisfying condition are not included in query.

Iris_2d Domain

- iris_2d (petallength float, petalwidth float, class varchar(30));
- Average petal length and width for Iris Setosa, Iris Versicolor, Iris Virginica
 - SELECT class, avg(petallength), avg(petalwidth)
FROM iris_2d
GROUP BY class
HAVING min(petallength) > 1;

class	avg(petallength)	avg(petalwidth)
Iris-versicolor	4.259999981	1.325999992
Iris-virginica	5.551999989	2.025999978

Having Clause

- Same Constraints as Select w/ Group By
 - Anything goes w/ a subquery
 - Grouping Attributes w/ condition.
 - Grouping Operators w/ condition.

Iris Domain



- Iris Versicolor



- Iris Setosa



- Iris Virginica

Iris_2d

petallength	petalwidth	class
1.4	0.2	Iris-setosa
1.4	0.2	Iris-setosa
1.3	0.2	Iris-setosa
1.5	0.2	Iris-setosa
1.4	0.2	Iris-setosa
3.9	1.1	Iris-versicolor
4.8	1.8	Iris-versicolor
4	1.3	Iris-versicolor
4.9	1.5	Iris-versicolor
4.7	1.2	Iris-versicolor
5.2	2.3	Iris-virginica
5	1.9	Iris-virginica
5.2	2	Iris-virginica
5.4	2.3	Iris-virginica
5.1	1.8	Iris-virginica

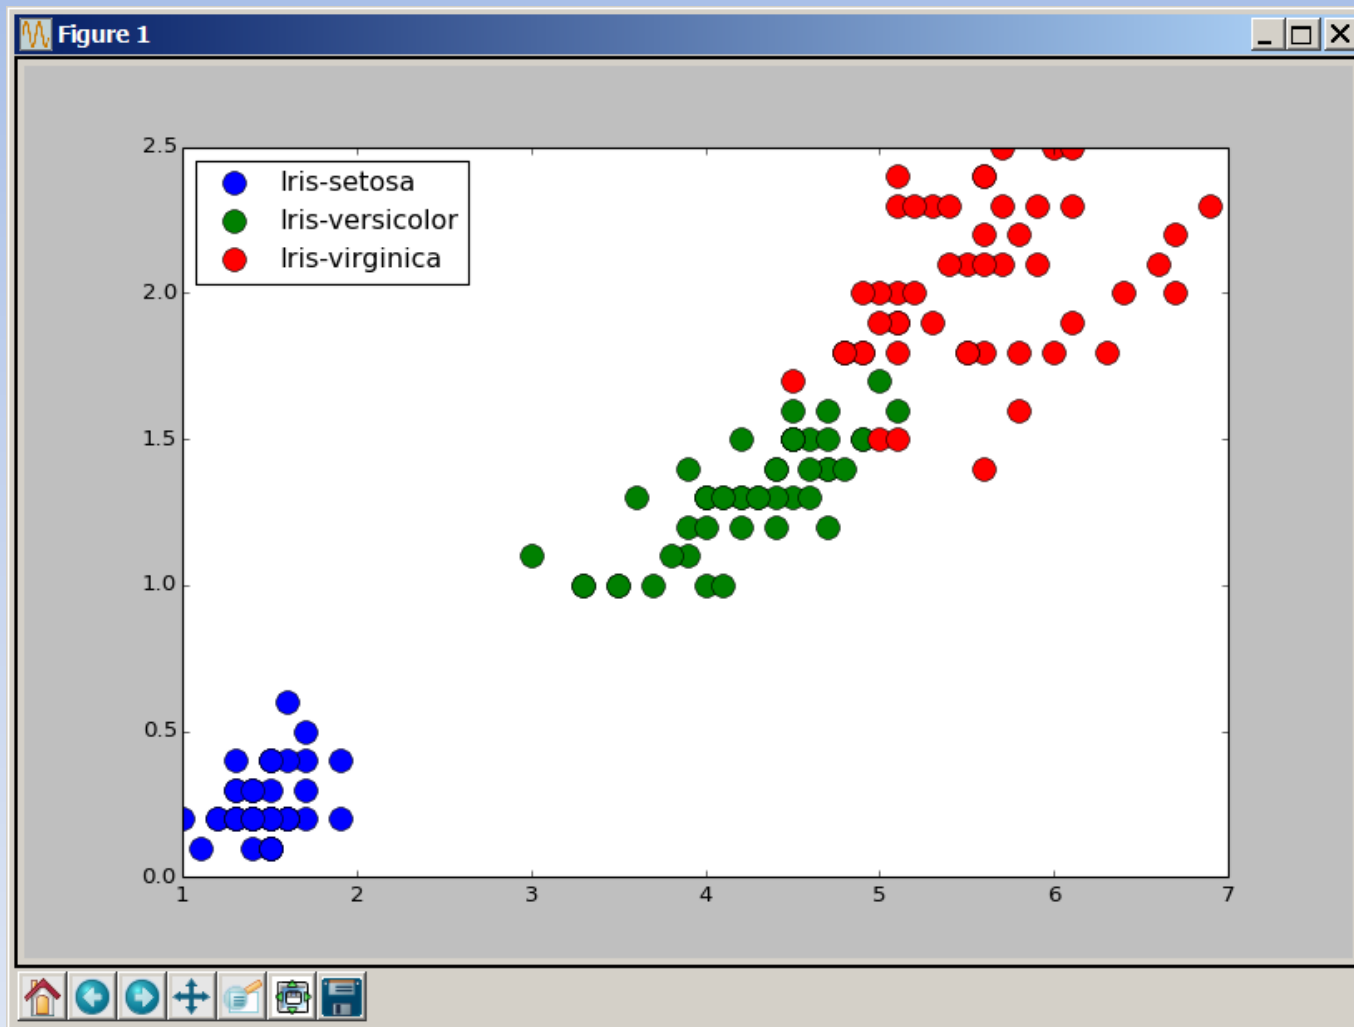
Iris_2d Domain

```
def plot(db, table, x_item, y_item, ciris):
    conn = mysql.connector.connect (host = "localhost", user="root", passwd = "cs126", db = db)
    cursor = conn.cursor ()
    s = "select " +x_item+", "+y_item+" from " + table \
        + " where class like '"+ciris+"%"
    cursor.execute (s)
    row = cursor.fetchone ()
    X = y= []
    while row != None:
        X.append(row[0])
        Y.append(row[1])
        row = cursor.fetchone ()
    cursor.close ()
    conn.close ()
    matplotlib.pyplot.plot(X,Y,'o', markersize=12, label= ciris)
```

Iris_2d Domain

```
def test2():  
    db = "iris2d"  
    table = "iris_2d"  
    x_item = 'petallength'  
    y_item = 'petalwidth'  
    plot(db, table, x_item, y_item, 'Iris-setosa')  
    plot(db, table, x_item, y_item, 'Iris-versicolor')  
    plot(db, table, x_item, y_item, 'Iris-virginica')  
    matplotlib.pyplot.legend(loc='upper left', numpoints = 1)  
    matplotlib.pyplot.show()
```

Iris_2d Domain



Questions 6.4.6

- Product(maker, model, type)
- PC(model, speed, ram, hd, price)
- Laptop(model, speed, ram, hd, screen, price)
- Printer(model, color, type, price)
- a) Find the average speed of PC's .

Questions 6.4.6

```
698 • use computers;  
699 • select * from product;  
700 • select avg(speed) as avg_speed from pc;  
701
```

Result Set Filter:



Export:



Wrap Cell Content:



avg_speed

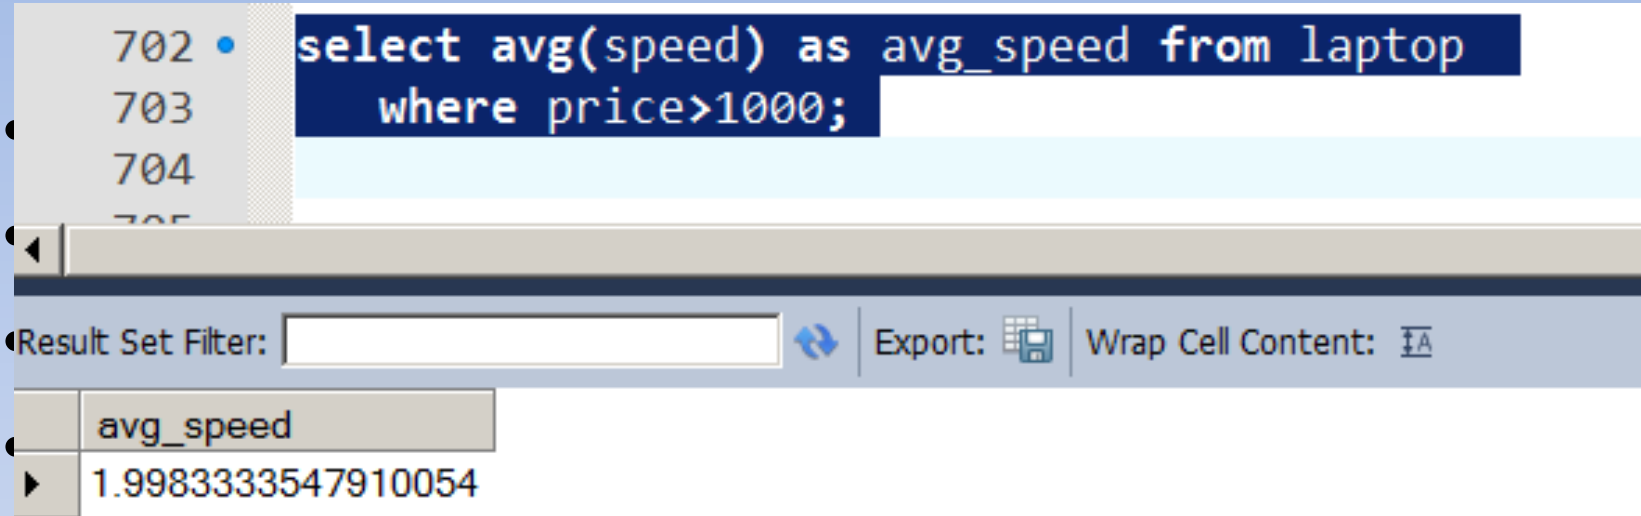
2.4846153809474063

select avg(speed) as avg_speed from pc;

Questions 6.4.6

- Product(maker, model, type)
- PC(model, speed, ram, hd, price)
- Laptop(model, speed, ram, hd, screen, price)
- Printer(model, color, type, price)
- b) Find the average speed of laptops costing over \$1000.



Questions 6.4.6



The screenshot shows a SQL query editor with the following text:

```
702 • select avg(speed) as avg_speed from laptop
703     where price>1000;
704
705
```

Below the editor is a toolbar with the following options:

- Result Set Filter:
- Export: 
- Wrap Cell Content: 

Below the toolbar is a table with the following data:

avg_speed
1.9983333547910054

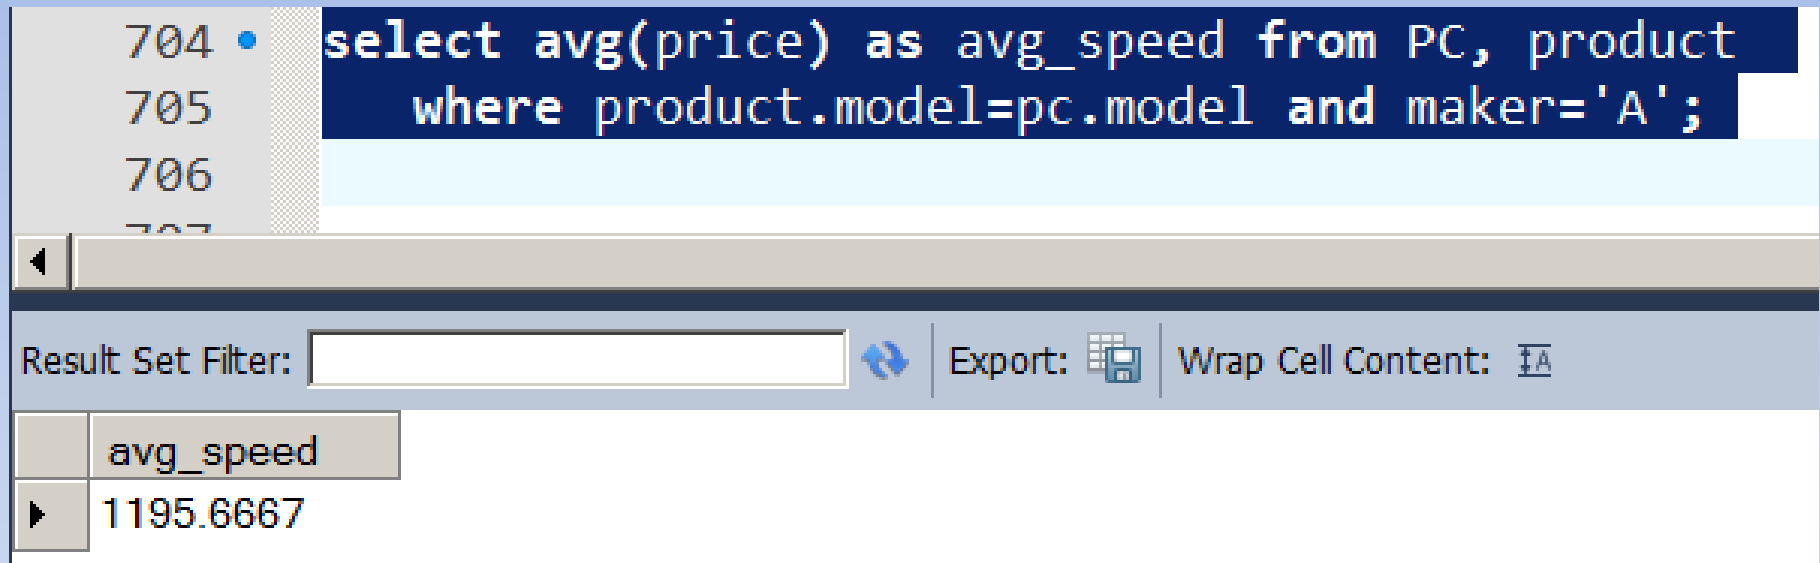
- b) Find the average speed of laptops costing over \$1000.

```
select avg(speed) as avg_speed from laptop
where price>1000;
```

Questions 6.4.6


- Product(maker, model, type)
- PC(model, speed, ram, hd, price)
- Laptop(model, speed, ram, hd, screen, price)
- Printer(model, color, type, price)
- c) Find the average price of PC's made by manufacturer "A ."

Questions 6.4.6



The screenshot shows a SQL query editor with a dark blue background. The query is: `select avg(price) as avg_speed from PC, product where product.model=pc.model and maker='A';`. Below the query editor is a result set table with a light blue header and a white body. The header row has a column named `avg_speed`. The first data row shows the value `1195.6667`. The interface includes a 'Result Set Filter' input field, an 'Export' button with a floppy disk icon, and a 'Wrap Cell Content' checkbox.

```
704 • select avg(price) as avg_speed from PC, product
705     where product.model=pc.model and maker='A';
706
707
```

Result Set Filter: Export:  Wrap Cell Content: ☐

avg_speed
1195.6667

c) Find the average price of PCs made by manufacturer "A."

```
select avg(price) as avg_price from PC, product
where product.model=pc.model and maker='A';
```

Questions 6.4.6

- Product(maker, model, type)
- PC(model, speed, ram, hd, price)
- Laptop(model, speed, ram, hd, screen, price)
- Printer(model, color, type, price)
- c) Find , for each different speed, the average price of a PC .

```

707 • SELECT  SPEED,
708          AVG(price) AS AVG_PRICE
709 FROM    PC
710 GROUP BY speed ;
711
712

```

- Product
- PC(n)
- Laptop
- Printer
- c) Find the average price of

Result Set Filter: Export: Wrap Cells

	SPEED	AVG_PRICE
▶	1.42	478.0000
	1.86	959.0000
	2	650.0000
	2.1	995.0000
	2.2	640.0000
	2.66	2114.0000
	2.8	689.3333
	3.06	529.0000
	3.2	839.5000

Result 8 x

select speed, avg(price) as avg_price from PC
group by speed ;

Questions 6.4.6

- Product(maker, model, type)
 - PC(model, speed, ram, hd, price)
 - Laptop(model, speed, ram, hd, screen, price)
 - Printer(model, color, type, price)
-
- g) Find the manufacturers that make at least three different models of PC .

- Product(mak
- PC(model, sp
- Laptop(mod
- Printer(mod

```

712 • SELECT R.make
713 FROM Product R,
714 PC P
715 WHERE R.model = P.model
716 GROUP BY R.make
717 HAVING COUNT(R.model) >=3 ;
718
719

```

Result Set Filter: Export:

maker
A
B
D
E

- g) Find the maker of PC. different models

```

SELECT R.make
FROM Product R,
      PC P
WHERE R.model = P.model
GROUP BY R.make
HAVING COUNT(R.model) >=3 ;


```

```

719 • SELECT maker
720 FROM Product
721 WHERE ctype='pc'
722 GROUP BY maker
723 HAVING COUNT(model) >=3 ;
724

```

- Product(ma
- PC(model, s
- Laptop(mod
- Printer(mod

Result Set Filter: Export: 

maker
A
B
D
E

- g) Find the manufacturers that make at least three different models of PC.

```

SELECT maker
FROM Product
WHERE ctype='pc'
GROUP BY maker
HAVING COUNT(model) >=3 ;

```


In-Class


- h) Find for each manufacturer who sells PC's the maximum price of a PC.

```

725 • SELECT R.maker,
726         MAX(P.price) AS Max_Price
727 FROM   Product R,
728        PC P
729 WHERE  R.model = P.model
730 GROUP BY R.maker ;
731
732

```

- The result set shows the maximum price for each maker.

Result Set Filter: Export:  Wrap Cell Contents

	maker	Max_Price
FROM	A	2114
	B	1049
	C	510
WHERE	D	770
	E	959
GROUP BY R.maker ;		