

### Indexes

- Index = data structure used to speed access to tuples of a relation, given values of one or more attributes.
- Could be a hash table, but in a DBMS it is always a balanced search tree with giant nodes (a full disk page) called a *B-tree*.



### Recursion in SQL

Databases - DB14 Started - Jun 09, 2014



View Course

Your final grade: 100%.

Download Statement (PDF)



### On-Line Analytical Processing

Databases - DB13 Started - Jun 09, 2014



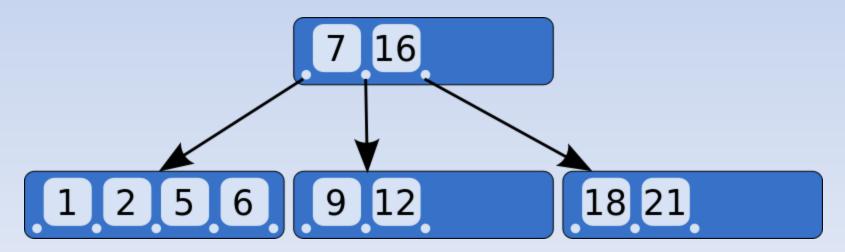
View Course

Your final grade: 100%.

Download Statement (PDF)

### **B-Tree: Quick Glance**

- According to Knuth's definition, a B-tree of order m is a tree which satisfies the following properties:
  - Every node has at most m children.
  - Every non-leaf node (except root) has at least  $\lceil m/2 \rceil$  children.
  - The root has at least two children if it is not a leaf node.
  - A non-leaf node with k children contains k-1 keys.
  - All leaves appear in the same level.
- For example, in a 2-3 B-Tree (often simply referred to as a 2-3 tree), each internal node may have only 2 or 3 child nodes.
- Example below is 3-5 Tree



## **Declaring Indexes**

- No standard!
- Typical syntax:

```
CREATE INDEX PersonInd ON
  Person(Pid);
CREATE INDEX MovieInd ON
  Movies(mid);
```

- Given a value v, the index takes us to only those tuples that have v in the attribute(s) of the index.
- Example: use PersonInd and MovieInd to find ratings for person and movie.

- Given a value v, the index takes us to only those tuples that have v in the attribute(s) of the index.
- Example: In Movies DB
  - use pid (Person ID) and mid (Movie ID) to find ratings for person and movie.

```
alter table movie_ratings add index (pid); alter table movie_ratings add index (mid);
```

- Given a value v, the index takes us to only those tuples that have v in the attribute(s) of the index.
- Example: In Movies DB
  - use pid (Person ID) and mid (Movie ID) to find ratings for person and movie.

```
alter table movie_ratings add index (pid); alter table movie_ratings add index (mid);
```

- Given a value v, the index takes us to only those tuples that have v in the attribute(s) of the index.
- Example: In Movies DB
  - use pid (Person ID) and mid (Movie ID) to find ratings for person and movie.

SELECT count(\*) FROM ratings a, ratings b WHERE a.mid = b.mid AND a.rating = b.rating AND a.pid < b.pid;

40 seconds w/ Index 10 minutes plus w/o Index

## **Database Tuning**

- A major problem in making a database run fast is deciding which indexes to create.
- Pro: An index speeds up queries that can use it.
- Con: An index slows down all modifications on its relation because the index must be modified too.

## **Example:** Tuning

- Suppose the only things we did with our Movie Ratings database was:
  - 1. Insert new facts into a relation (10%).
  - 2. Find the average rating for a movie (90%).

## **Tuning Advisors**

- A major research thrust.
  - Because hand tuning is so hard.
- An advisor gets a query load, e.g.:
  - 1. Choose random queries from the history of queries run on the database, or
  - 2. Designer provides a sample workload.

## Tuning Advisors --- (2)

- The advisor generates candidate indexes and evaluates each on the workload.
  - Feed each sample query to the query optimizer,
     which assumes only this one index is available.
  - Measure the improvement/degradation in the average running time of the queries.

### Question

#### MULTIPLE CHOICE (2/2 points)

[Q1] Consider the following relational schema:

```
Course(courseName unique, department, instrID)
Instructor(instrID unique, office)
Student(studentID unique, major)
Enroll(studentID, courseName, unique (studentID,courseName))
```

Suppose there are five types of queries commonly asked on this schema:

- Given a course name, find the department offering that course.
- List all studentIDs together with all of the departments they are taking courses in.
- Given a studentID, find the names of all courses the student is enrolled in.
- List the offices of instructors teaching at least one course.
- Given a major, return the studentIDs of students in that major.

Which of the following indexes could NOT be useful in speeding up execution of one or more of the above queries?

Index on Student.studentID

Index on Enroll.studentID

Index on Student.major

Index on Enroll.courseName

### Question

#### MULTIPLE CHOICE (2/2 points)

[Q1] Consider the following relational schema:

```
Course (courseName unique, department, instrID)
Instructor(instrID unique, office)
Student(studentID unique, major)
Enroll(studentID, courseName, unique (studentID, courseName))
```

Suppose there are five types of queries commonly asked on this schema:

- Given a course name, find the department offering that course.
- List all studentIDs together with all of the departments they are taking courses in.
- Given a studentID, find the names of all courses the student is enrolled in.
- List the offices of instructors teaching at least one course.
- Given a major, return the studentIDs of students in that major.

Which of the following indexes could NOT be useful in speeding up execution of one or more of the above queries?

- Index on Student.studentID
- C Index on Enroll.studentID
- Index on Student.major
- C Index on Enroll.courseName

[Q2] Consider a table storing temperature readings taken by sensors:

```
Temps(sensorID, time, temp)
```

Assume the pair of attributes [sensorID,time] is a key. Consider the following query:

```
select * from Temps
where sensorID = 'sensor541'
and time = '05:11:02'
```

Consider the following scenarios:

- A No index is present on any attribute of Temps
- B An index is present on attribute sensorID only
- C An index is present on attribute time only
- D Separate indexes are present on attributes sensorID and time
- E A multi-attribute index is present on (sensorID,time)

Suppose table Temps has 50 unique sensorIDs and each sensorID has exactly 20 readings. Furthermore there are exactly 10 readings for every unique time in Temps.

For each scenario A-E, determine the maximum number of tuples that might be accessed to answer the query, assuming one "best" index is used whenever possible. (Don't count the number of index accesses.) Which of the following combinations of values is correct?

```
A:1000, B:20, E:1
A:1000, D:20, E:10
C:1000, D:20, E:1
A:1000, B:20, E:10
```

[Q2] Consider a table storing temperature readings taken by sensors:

```
Temps(sensorID, time, temp)
```

Assume the pair of attributes [sensorID,time] is a key. Consider the following query:

```
select * from Temps
where sensorID = 'sensor541'
and time = '05:11:02'
```

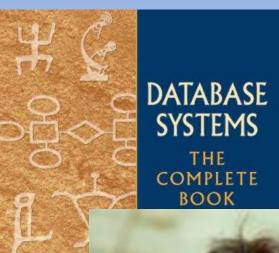
Consider the following scenarios:

- A No index is present on any attribute of Temps
- B An index is present on attribute sensorID only
- C An index is present on attribute time only
- D Separate indexes are present on attributes sensorID and time
- E A multi-attribute index is present on (sensorID,time)

Suppose table Temps has 50 unique sensorIDs and each sensorID has exactly 20 readings. Furthermore there are exactly 10 readings for every unique time in Temps.

For each scenario A-E, determine the maximum number of tuples that might be accessed to answer the query, assuming one "best" index is used whenever possible. (Don't count the number of index accesses.) Which of the following combinations of values is correct?

- A:1000, B:20, E:1
- O A:1000, D:20, E:10
- C:1000, D:20, E:1
- <sup>C</sup> A:1000, B:20, E:10



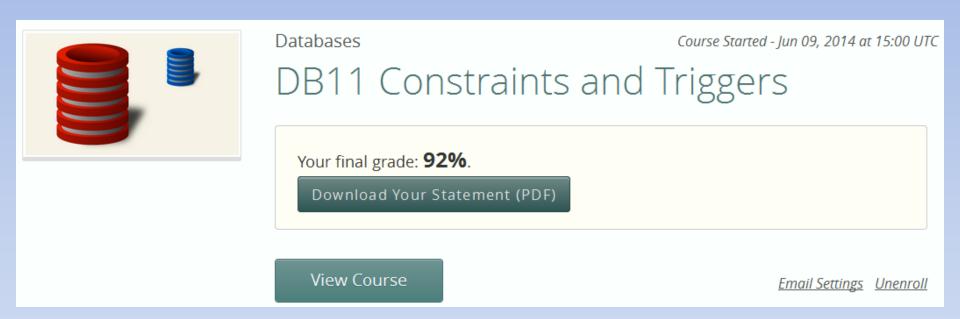
**Chapter 7:** Constraints and **Triggers** 

Chapter 9: SQL in a Server Environment

**Stored Procedures** 



### Stanford Online



#### **Triggers**

"Event-Condition-Action Rules"

When event occurs, check condition; if true, do action

- 1) Move monitoring logic from apps into DBMS
- 2) Enforce constraints
  - Beyond what constraint system supports
  - Automatic constraint "repair"
- \* Implementations vary significantly

#### **Triggers in SQL**

```
Create Trigger name
Before|After|Instead Of [insert|delete|update]

[For Each Row]

When (condition)

action
```

## Triggers in SQL Calculate the total amount inserted.

```
CREATE TRIGGER tr_Accountsum

AFTER INSERT ON BankAccount

FOR EACH ROW

BEGIN

set @sum = @sum + new.amount;

END
```

# Triggers in SQL Calculate the total amount inserted.

```
set @sum=0;
insert into BankAccount values
(123, 801, 500),
(235, 701, 500),
(235, 702, 500),
select @sum;
```

## Triggers in SQL Calculate the total amount inserted.

```
87
     set @sum=0;
88
     insert into BankAccount values
89 •
         (100, 901, 500),
90
         (123, 801, 500),
91
         (235, 701, 500),
92
         (235, 702, 500),
93
         (400, 401, 500);
94
     select @sum;
95
96
     select trigger schema, Trigger name, Event manipulation
97 •
         from information_schema.Triggers;
98
99
                   💎 File: 🍱 Autosize: 🏗
@sum
2500
```

#### **Triggers in SQL**

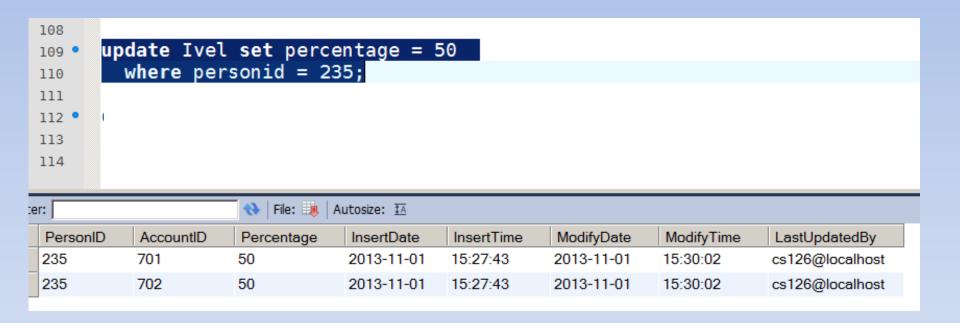
#### Monies distribution between accounts table.

```
CREATE TABLE Ivel(
 PersonID char(3), AccountID int,
 Percentage int,
InsertDate date default null,
InsertTime time default null,
ModifyDate date default null,
ModifyTime time default null,
LastUpdatedBy varchar(132) default null,
FOREIGN KEY (PersonID) REFERENCES
  person(PersonID)
  On Delete Cascade On Update Set Null,
FOREIGN KEY (AccountID) REFERENCES
  BankAccount(AccountID)
   On Delete Cascade On Update Set Null);
```

```
CREATE TRIGGER tr_IvelInsertDate
   before insert on Ivel
for each row
begin
   set new.InsertDate = curdate();
   set new.InsertTime = curtime();
   set new.ModifyDate = curdate();
   set new.ModifyTime = curtime();
   set new.LastUpdatedBy = user();
end
```

```
CREATE TRIGGER tr_IvelModifyDate
before update on Ivel
for each row
begin
  set new.ModifyDate = curdate();
  set new.ModifyTime = curtime();
  set new.LastUpdatedBy = user();
end
```

```
insert into Ivel(PersonID, AccountID, Percentage) values
 103 •
            (235, 701, 75),
 104
            (235, 702, 25);
 105
 106
         select * from Ivel;
 107 •
 108
ter:
                          💎 | File: 🍱 | Autosize: 🏗
             AccountID
  PersonID
                          Percentage
                                       InsertDate
                                                   InsertTime
                                                                ModifyDate
                                                                             ModifyTime
                                                                                          LastUpdatedBy
  235
                                      2013-11-01
                                                   15:27:43
                                                               2013-11-01
                                                                            15:27:43
                                                                                         cs126@localhost
             701
                         75
  235
             702
                         25
                                      2013-11-01
                                                   15:27:43
                                                               2013-11-01
                                                                            15:27:43
                                                                                         cs126@localhost
```



# Triggers in SQL Inserting an account where person doesn't exist.

```
CREATE TRIGGER tr_NewPerson
before insert on BankAccount
for each row
begin
if not exists
  (select * from Person
    where PersonID = new.PersonID) then
  set new.PersonID = 000;
end if;
end
```

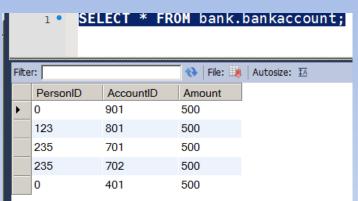
# Triggers in SQL Inserting an account where person doesn't exist.

```
insert into Person
  values(000, NULL, '0000-00-00');
set @sum=0;
insert into BankAccount values
  (100, 901, 500), (123, 801, 500),
  (235, 701, 500), (235, 702, 500),
  (400, 401, 500);
select @sum;
```

#### **Triggers in SQL**

#### Inserting an account where person doesn't exist.

```
insert into Person
          values(123, 'John Sigma123', '1933-01-01');
   77
       insert into Person
         values(235, 'John Sigma235', '1955-01-01');
  79
       insert into Person
   80 •
         values(000, NULL, '0000-00-00');
   81
   82
   83
       delete from BankAccount;
   84 •
       insert into BankAccount values(100, 901, 500);
  se • insert into BankAccount values(123, 801, 500);
  87 • insert into BankAccount values (235, 701, 500);
       insert into BankAccount values(235, 702, 500);
       insert into BankAccount values(400, 401, 500);
   90
       set @sum=0;
       insert into BankAccount values
          (100, 901, 500),
        (123, 801, 500),
        (235, 701, 500),
        (235, 702, 500),
         (400, 401, 500);
       select @sum;
Filter:
                     💎 | File: 🍱 | Autosize: 🏗
  @sum
  2500
```



#### **Triggers in SQL (general syntax)**

```
Create Trigger name
Before | After | Instead Of events
[ referencing-variables ]
[For Each Row]
When (condition)
action
```

### Exercise 7.2.2a

Write the following constraints on attribues from our example schema:

Product(maker, model, ctype)
PC(model, speed, ram, hd, price)
Laptop(model, speed, ram, hd, screen, price)
Printer(model, color, type, price)

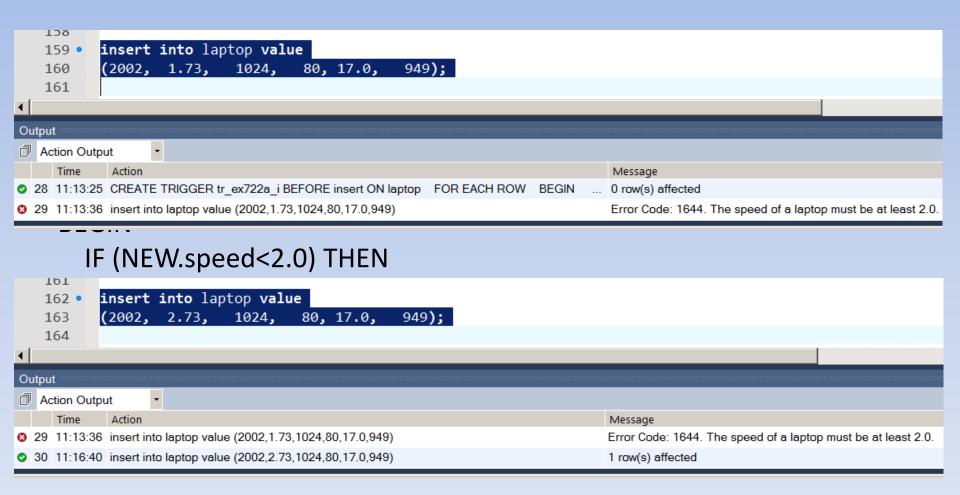
a) The speed of a laptop must be at least 2.0.

In MySQL we'll do this with a Trigger.

### Exercise 7.2.2a

```
The speed of a laptop must be at least 2.0.
delimiter //
CREATE TRIGGER tr_ex722a_i
BEFORE insert ON laptop
FOR EACH ROW
BEGIN
  IF (NEW.speed<2.0) THEN
    SIGNAL sqlstate '45000'
     SET message_text = 'The speed of a laptop must be at least 2.0.';
  END IF;
END//
delimiter;
```

### Exercise 7.2.2a



# Exercise 7.2.2a A little more advanced.

```
The speed of a laptop must be at least 2.0.
delimiter //
CREATE TRIGGER tr ex722a i
BEFORE insert ON laptop
FOR FACH ROW
BEGIN
    DECLARE msg varchar(255);
    IF (NEW.speed<2.0)
    THFN
       SET msg =
         concat('Constraint tr_ex722a_i violated: Laptop speed must be >2.0: ',
          cast(new.speed as char));
       SIGNAL sqlstate '45000' SET message text = msg;
    END IF;
  END//
delimiter;
```

### Exercise 7.2.2a

The speed of a laptop must be at least 2.0. delimiter // CREATE TRIGGER tr ex722a i **BEFORE** insert ON laptop **FOR FACH ROW** delimiter // 116 CREATE TRIGGER tr ex722a i BEFORE insert ON laptop 117 • FOR EACH ROW 118 BEGIN 119 DECLARE msg varchar(255); 120 IF (NEW.speed<2.0)</pre> 121 THEN 122 SET msg = concat('Constraint tr\_ex722a\_i violated: Laptop speed must be >2.0: ', cast(new.speed as char)); 123 SIGNAL sqlstate '45000' SET message text = msg; 124 END IF; 125 END// 126 delimiter ; 127 128 Output Action Output Time Action Message 18 11:01:48 drop trigger tr\_ex722a\_i 0 row(s) affected

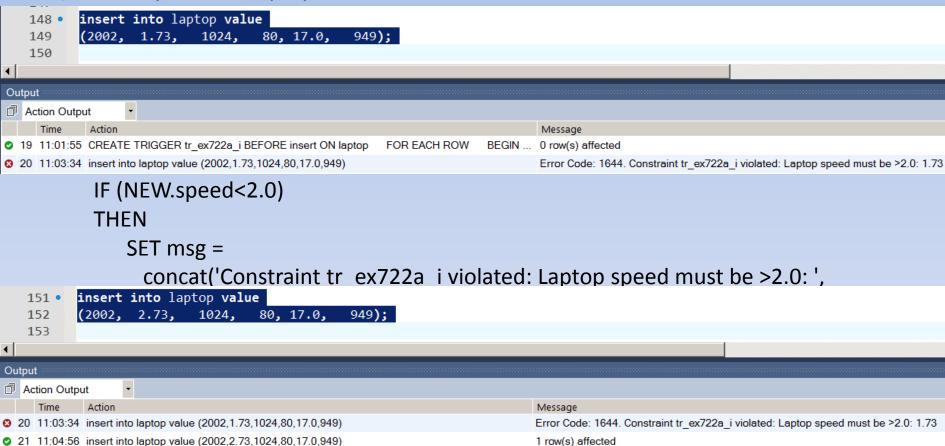
BEGIN ... 0 row(s) affected

FOR EACH ROW

19 11:01:55 CREATE TRIGGER tr\_ex722a\_i BEFORE insert ON laptop

### Exercise 7.2.2a

a) The speed of a laptop must be at least 2.0.



# Exercise 7.2.2a Will Need: Update & Insert Trigger

```
The speed of a laptop must be at least 2.0.
delimiter //
CREATE TRIGGER tr ex722a u
BEFORE update ON laptop
  FOR FACH ROW
  BEGIN
    DECLARE msg varchar(255);
    IF (NEW.speed<2.0)
    THEN
     SET msg =
      concat('Constraint tr_ex722a_u violated: Laptop speed must be >2.0: ',
      cast(new.speed as char));
     SIGNAL sqlstate '45000' SET message text = msg;
    END IF:
  END//delimiter;
```

# Exercise 7.2.2a Will Need: Update & Insert Trigger

```
The speed of a laptop must be at least 2.0.
   delimiter //
   CREATE TRIGGER tr ex722a u
   BEFORE update ON laptop
       FOR FACH ROW
       BEGIN
         DECLARE msg varchar(255);
         IF (NEW.speed<2.0)
         THEN
           SFT msg =
   164
         update laptop set speed = 1.73 where model=2002;
   165
Output
Action Output
                                                                       Message
34 11:22:57 CREATE TRIGGER tr_ex722a_u BEFORE update ON laptop
                                                    FOR EACH ROW
                                                                 BEGI... 0 row(s) affected
 35 11:23:02 update laptop set speed = 1.73 where model=2002
                                                                       Error Code: 1644. Constraint tr ex722a u violated: Laptop speed must be >2.0: 1.73
```

 Write the following as triggers. In ease case, disallow or undo the modification if it does not satisfy the stated constraint. The database schema is from the battleships schema.

```
Classes(class, stype, country, numGuns, bore, displacement);
Battles(bname, bdate);
Outcomes(ship, battle, result);
Ships(sname, class, launched);
```

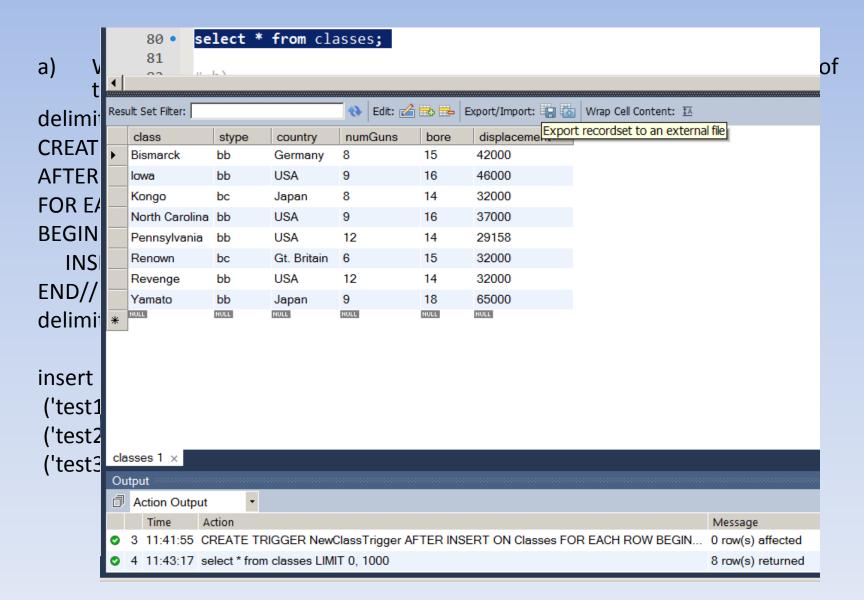
a) When a new class is inserted into Classes, also insert a ship with the name of that class and a NULL launch date.

a) When a new class is inserted into Classes, also insert a ship with the name of that class and a NULL launch date.

```
delimiter //
CREATE TRIGGER NewClassTrigger
AFTER INSERT ON Classes
FOR EACH ROW
BEGIN
INSERT INTO Ships VALUE (new.class, New.class, NULL);
END//
delimiter;
```

a) When a new class is inserted into Classes, also insert a ship with the name of that class and a NULL launch date.

```
delimiter //
CREATE TRIGGER NewClassTrigger
AFTER INSERT ON Classes
FOR EACH ROW
BEGIN
  INSERT INTO Ships VALUE (new.class, New.class, NULL);
END//
delimiter;
insert into classes values
('test1', 'bb', 'USA', 8, 15, 33000),
('test2', 'bb', 'USA', 8, 15, 36000),
('test3', 'bb', 'USA', 8, 15, 42000);
```



```
insert into classes values
('test1', 'bb', 'USA', 8, 15, 33000),
('test2', 'bb', 'USA', 8, 15, 36000),
('test3', 'bb', 'USA', 8, 15, 42000);

Output

Time Action
Action Output

Time Action
Action
Action

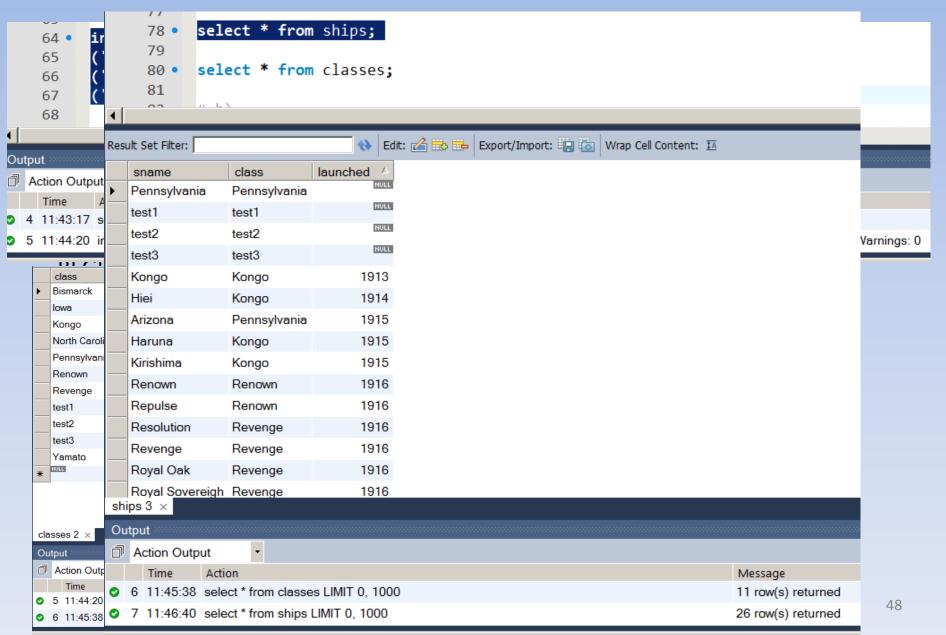
Message

1 11:43:17 select * from classes LIMIT 0, 1000

5 11:44:20 insert into classes values ('test1', 'bb', 'USA', 8, 15, 33000), ('test2', 'bb', 'USA', 8, 15, 360... 3 row(s) affected Records: 3 Duplicates: 0 Warnings: 0
```

	DL7'INI					
	class	stype	country	numGuns	bore	displacement
١	Bismarck	bb	Germany	8	15	42000
	Iowa	bb	USA	9	16	46000
	Kongo	bc	Japan	8	14	32000
	North Carolina	bb	USA USA	9	16	37000
	Pennsylvania	bb		12	14	29158
	Renown	bc	Gt. Britain	6	15	32000
	Revenge	bb	USA	12	14	32000
	test1	bb	USA	8	15	33000
	test2	bb	USA	8	15	35000
	test3	bb	USA	8	15	35000
	Yamato	bb	Japan	9	18	65000
*	NULL	NULL	NULL	NULL	NULL	NULL

cli	classes 2 ×				
0	Output				
ā	Action Out	out •			
	Time	Action	Message		
0	5 11:44:20	insert into classes values ('test1', 'bb', 'USA', 8, 15, 33000), ('test2', 'bb', 'USA', 8, 15, 360	3 row(s) affected Records: 3 Duplicates: 0 Warnings: 0		
0	6 11:45:38	select * from classes LIMIT 0, 1000	11 row(s) returned		



[Q2] Consider the following trigger over a table R(a,b), specified using the trigger language of the SQL standard:

```
CREATE TRIGGER Rins

AFTER INSERT ON R

REFERENCING NEW ROW AS new

FOR EACH ROW

INSERT INTO R(a,b)

(SELECT DISTINCT R.a, new.b

FROM R

WHERE R.b = new.a)

EXCEPT

(SELECT DISTINCT a,b FROM R)
```

Suppose table R is empty initially. We issue three commands to insert tuples into R: first we insert (1,2), then we insert (2,3), then we insert (3,4). After some of these inserts, trigger **Rins** may insert further tuples into R, which may activate the trigger recursively. After all the inserts are done, which of these tuples is NOT in table R?

- (2,4)
- (1,4)
- (3,1)
- (1,3)

[Q2] Consider the following trigger over a table R(a,b), specified using the trigger language of the SQL standard:

```
CREATE TRIGGER Rins

AFTER INSERT ON R

REFERENCING NEW ROW AS new

FOR EACH ROW

INSERT INTO R(a,b)

(SELECT DISTINCT R.a, new.b

FROM R

WHERE R.b = new.a)

EXCEPT

(SELECT DISTINCT a,b FROM R)
```

Suppose table R is empty initially. We issue three commands to insert tuples into R: first we insert (1,2), then we insert (2,3), then we insert (3,4). After some of these inserts, trigger **Rins** may insert further tuples into R, which may activate the trigger recursively. After all the inserts are done, which of these tuples is NOT in table R?

```
○ (2,4)
○ (1,4)
◎ (3,1) ✓
○ (1,3)
```

Students at your hometown high school have decided to organize their social network using databases. So far, they have collected information about sixteen students in four grades, 9-12. Here's the schema:

Highschooler (ID, name, grade)

English: There is a high school student with unique *ID* and a given *first name* in a certain *grade*.

Friend (ID1, ID2)

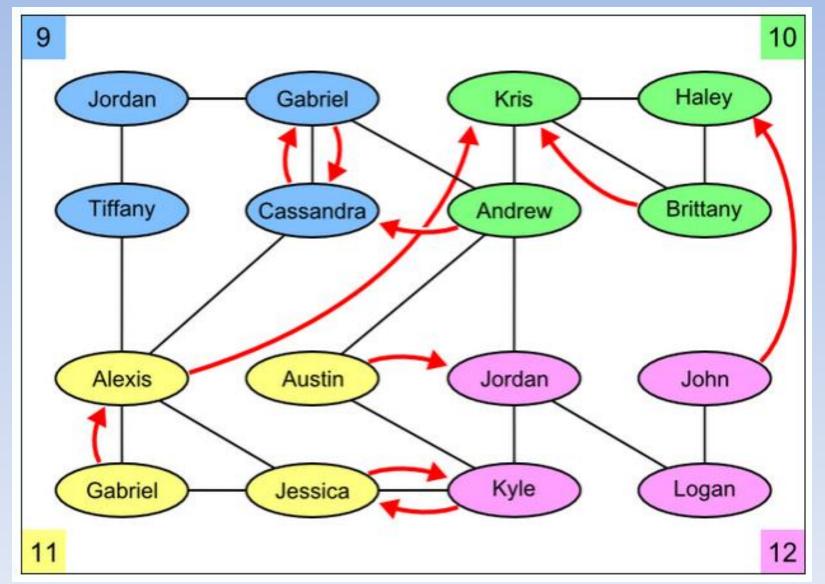
English: The student with *ID1* is friends with the student with *ID2*. Friendship is mutual, so if (123, 456) is in the Friend table, so is (456, 123).

Likes (ID1, ID2)

English: The student with *ID1* likes the student with *ID2*. Liking someone is not necessarily mutual, so if (123, 456) is in the Likes table, there is no guarantee that (456, 123) is also present.

Your triggers will run over a small data set conforming to the schema. View the database. (You can also download the schema and data.)

For your convenience, here is a graph showing the various connections between the people in our database. 9th graders are blue, 10th graders are green, 11th graders are yellow, and 12th graders are purple. Undirected black edges indicate friendships, and directed red edges indicate that one person likes another person.



Q1 (1/1 point)
Write a trigger that makes new students named 'Friendly' automatically like everyone else in their grade. That is, after the trigger runs, we should have ('Friendly', A) in the Likes table for every other Highschooler A in the same grade as 'Friendly'.
<ul> <li>Your triggers are created in SQLite, so you must conform to the trigger constructs supported by SQLite.</li> </ul>

#### Q1 (1/1 point)

Write a trigger that makes new students named 'Friendly' automatically like everyone else in their grade. That is, after the trigger runs, we should have ('Friendly', A) in the Likes table for every other Highschooler A in the same grade as 'Friendly'.

 Your triggers are created in SQLite, so you must conform to the trigger constructs supported by SQLite.

```
CREATE TRIGGER q1 after insert ON highschooler

FOR EACH ROW when new.name='Friendly'

BEGIN

insert into Likes select new.id, id from highschooler where id<>new.id and new.grade = grade;

END
```

To check your trigger(s), we first ran the following data modification statement(s): *insert into Highschooler values (1000, 'Friendly', 9)*;

insert into Highschooler values (2000, 'Friendly', 11); insert into Highschooler values (3000, 'Unfriendly', 10).

We then ran the following query:  $select\ H1.name,\ H1.grade,\ H2.name,\ H2.grade\ from\ Likes\ L,\ Highschooler\ H1,\ Highschooler\ H2$  where L.ID1 = H1.ID and L.ID2 = H2.ID order by  $H1.name,\ H1.grade,\ H2.name,\ H2.grade$ 

Alexis	11	Kris	10
Andrew	10	Cassandra	9
Austin	11	Jordan	12
Brittany	10	Kris	10
Cassandra	9	Gabriel	9
Friendly	9	Cassandra	9
Friendly	9	Gabriel	9
Friendly	9	Jordan	9
Friendly	9	Tiffany	9
Friendly	11	Alexis	11
Friendly	11	Austin	11
Friendly	11	Gabriel	11
Friendly	11	Jessica	11
Gabriel	9	Cassandra	9
Gabriel	11	Alexis	11
Jessica	11	Kyle	12
John	12	Haley	10
Kyle	12	Jessica	11

#### Q2 (1/1 point)

Write one or more triggers to manage the grade attribute of new Highschoolers. If the inserted tuple has a value less than 9 or greater than 12, change the value to NULL. On the other hand, if the inserted tuple has a null value for grade, change it to 9.

- Your triggers are created in SQLite, so you must conform to the trigger constructs supported by SQLite.
- To create more than one trigger, separate the triggers with a vertical bar (|).

- Two Parts:
  - Trigger Condition to find right situation.
  - Update Query to make the actual change.

- Trigger 1: "If the inserted tuple has a value less than 9 or greater than 12, change the value to NULL."
  - After Insert

CREATE TRIGGER q1 after insert ON highschooler

- Trigger 1: "If the inserted tuple has a value less than 9
  or greater than 12, change the value to NULL."
  - After Insert
  - For Each Row When...

```
CREATE TRIGGER q1 after insert ON highschooler

FOR EACH ROW when new.grade<9 or new.grade > 12

BEGIN

...make changes...

END
```

- Trigger 1: "If the inserted tuple has a value less than 9 or greater than 12, change the value to NULL."
  - After Insert
  - For Each Row When...
  - Update Query:

```
CREATE TRIGGER q1 after insert ON highschooler
FOR EACH ROW when new.grade<9 or new.grade > 12
BEGIN
update highschooler set grade = null where id=new.id;
END
```

- Trigger 2: "On the other hand, if the inserted tuple has a null value for grade, change it to 9."
  - After Insert
  - For Each Row When...
  - Update Query:

```
CREATE TRIGGER q1 after insert ON highschooler

FOR EACH ROW when new.grade<9 or new.grade > 12

BEGIN

update highschooler set grade = null where id=new.id;

END

CREATE TRIGGER q2 after insert ON highschooler

FOR EACH ROW when ______

BEGIN

update highschooler set ______ where _____;

END
```

#### Q2 (1/1 point)

Write one or more triggers to manage the grade attribute of new Highschoolers. If the inserted tuple has a value less than 9 or greater than 12, change the value to NULL. On the other hand, if the inserted tuple has a null value for grade, change it to 9.

- Your triggers are created in SQLite, so you must conform to the trigger constructs supported by SQLite.
- To create more than one trigger, separate the triggers with a vertical bar (|).

```
1 CREATE TRIGGER q1 after insert ON highschooler
 2
       FOR EACH ROW when new.grade<9 or new.grade > 12
 3
       BEGIN
 4
          update highschooler set grade = null where id=new.id;
 5
       END
6
7 CREATE TRIGGER q2 after insert ON highschooler
 8
       FOR EACH ROW when new grade is null
       BEGIN
10
          update highschooler set grade = 9 where id =new.id;
11
       END
12
```

To check your trigger(s), we first ran the following data modification statement(s): *insert into Highschooler* values (2121, 'Caitlin', null);

insert into Highschooler values (2122, 'Don', null);

insert into Highschooler values (2123, 'Elaine', 7);

insert into Highschooler values (2124, 'Frank', 20);

insert into Highschooler values (2125, 'Gale', 10)

We then ran the following query: select \* from Highschooler order by ID

1025	John	12
1101	Haley	10
1247	Alexis	11
1304	Jordan	12
1316	Austin	11
1381	Tiffany	9
1468	Kris	10
1501	Jessica	11
1510	Jordan	9
1641	Brittany	10
1661	Logan	12
1689	Gabriel	9
1709	Cassandra	9
1782	Andrew	10
1911	Gabriel	11
1934	Kyle	12
2121	Caitlin	9
2122	Don	9
2123	Elaine	<null></null>
2124	Frank	<null></null>
2125	Gale	10

# Views: Instead-Of Triggers



**Databases** 

Course Started - Jun 09, 2014 at 15:00 UTC

### DB12 Views and Authorization

Your final grade: **90%**.

Download Your Statement (PDF)

View Course

Email Settings Unenroll

You've started a new movie-rating website, and you've been collecting data on reviewers' ratings of various movies. Here's the schema:

**Movie** ( mID, title, year, director )

English: There is a movie with ID number *mID*, a *title*, a release *year*, and a *director*.

Reviewer ( rID, name )

English: The reviewer with ID number *rID* has a certain *name*.

**Rating** ( rID, mID, stars, ratingDate )

English: The reviewer *rID* gave the movie *mID* a number of *stars* rating (1-5) on a certain *ratingDate*.

In addition to the base tables, you've created three views:

View **LateRating** contains movie ratings after January 20, 2011. The view contains the movie ID, movie title, number of stars, and rating date.

create view LateRating as select distinct R.mID, title, stars, ratingDate from Rating R, Movie M where R.mID = M.mID and ratingDate > '2011-01-20'

View **HighlyRated** contains movies with at least one rating above 3 stars. The view contains the movie ID and movie title.

create view HighlyRated as select mID, title from Movie where mID in (select mID from Rating where stars > 3)

View **NoRating** contains movies with no ratings in the database. The view contains the movie ID and movie title.

create view NoRating as select mID, title from Movie where mID not in (select mID from Rating)

Q1 (1/1 point)					
Write an instead-of trigger that enables updates to the title attribute of view <b>LateRating</b> .					
<b>Policy:</b> Updates to attribute title in LateRating should update Movie.title for the corresponding movie. (You may assume attribute mID is a key for table Movie.) Make sure the mID attribute of view LateRating has not also been updated if it has been updated, don't make any changes. Don't worry about updates to stars or ratingDate.					
• Remember you need to use an instead-of trigger for view modifications as supported by SQLite.					

- Trigger Condition: Write an instead-of trigger that enables updates to the stars attribute of view LateRating.
  - Updates to attribute stars in LateRating should update Rating.stars for the corresponding movie rating.
    - You may assume attributes [mID,ratingDate] together are a key for table Rating.
  - Make sure the mID and ratingDate attributes of view LateRating have not also been updated
    - if either one has been updated, don't make any changes.
    - Don't worry about updates to title.

- Trigger: Write an instead-of trigger that enables updates to the stars attribute of view LateRating.
  - Instead Of Update on View
  - New Update Query on base tables

```
create trigger t1 instead of update of stars on LateRating
begin
update _____ set ___ = ____ where ____
end
```

#### Q1 (1/1 point)

Write an instead-of trigger that enables updates to the title attribute of view **LateRating**.

**Policy:** Updates to attribute title in LateRating should update Movie.title for the corresponding movie. (You may assume attribute mID is a key for table Movie.) Make sure the mID attribute of view LateRating has not also been updated -- if it has been updated, don't make any changes. Don't worry about updates to stars or ratingDate.

• Remember you need to use an instead-of trigger for view modifications as supported by SQLite.

```
create trigger t instead of update of title on laterating
begin

update movie set title = new.title where mid = new.mid;

end
```

To check your trigger, we first ran the following data modification statement(s): *insert into NoRating values* (104, 'E.T.'); *insert into NoRating values* (110, 'Avatar').

We then queried the view: select \* from NoRating View Result:

102	Star Wars
104	E.T.
105	Titanic

We then ran the following query: *select \* from Rating order by mID* Your Query Result:

201	101	2	2011-01-22
201	101	4	2011-01-27
202	106	4	<null></null>
203	103	2	2011-01-20
203	108	2	2011-01-30
203	108	4	2011-01-12
204	101	3	2011-01-09
205	103	3	2011-01-27
205	108	4	<null></null>
206	106	5	2011-01-19
206	107	3	2011-01-15
207	107	5	2011-01-20